

# ADVANCED SQL I

---

*CS 564- Fall 2018*

---

*ACKs: Dan Suciu, Jignesh Patel, AnHai Doan*

---

# WHAT IS THIS LECTURE ABOUT

---

- SQL: Set Operators
  - UNION/EXCEPT/INTERSECT
  - duplicates in SQL
- SQL: Nested Queries
  - IN/EXISTS/ALL
  - correlated queries

---

# SET AND MULTISSET OPERATORS

---

# SET OPERATORS: REFRESHER

$$R = \{1, 2, 3\}$$

$$S = \{1, 2, 4, 5\}$$

- Intersection:  $R \cap S = \{1, 2\}$
- Union:  $R \cup S = \{1, 2, 3, 4, 5\}$
- Difference:  
 $R - S = \{3\}$   
 $S - R = \{4, 5\}$

---

# SET OPERATORS IN SQL

---

SQL supports set operations between the outputs of subqueries:

- (subquery) **INTERSECT** (subquery)
- (subquery) **UNION** (subquery)
- (subquery) **EXCEPT** (subquery)

# SET OPERATORS: INTERSECT

SELECT A FROM R

**INTERSECT**

SELECT A FROM S;

R	A	S	A	output	A
	1		1		1
	1		1		
	1		2		
	2		2		2
	3		4		
			5		

Returns the tuples that belong in **both** subquery results

# SET OPERATORS: UNION

SELECT A FROM R

**UNION**

SELECT A FROM S;

R	A	S	A	output	A
	1		1		1
	1		1		2
	1		2		3
	2		2		4
	3		4		5
			5		

Returns the tuples that belong in **either** subquery results

# SET OPERATORS: EXCEPT

SELECT A FROM R  
**EXCEPT**  
SELECT A FROM S;

R	A	S	A	output	A
	1		1		3
	1		1		
	1		2		
	2		2		
	3		4		
			5		

Returns the tuples that belong in the first and **not** the second subquery result



# SEMANTICS

---

- When using set operators, SQL eliminates all duplicate tuples
- We can modify the semantics by using the keyword **ALL** (e.g. **UNION ALL**)
- When using **ALL**, the operators are evaluated using **multiset** (or **bag**) semantics

# SET OPERATORS: UNION ALL

SELECT A FROM R  
**UNION ALL**  
SELECT A FROM S;

output

R	A	S	A
	1		1
	1		1
	1		2
	2		2
	3		4
			5

The number of copies of each tuple is the **sum** of the number of copies in the subqueries

A
1
1
1
1
1
2
2
2
3
4
5

# SET OPERATORS: INTERSECT ALL

SELECT A FROM R  
**INTERSECT ALL**  
SELECT A FROM S;

R	A	S	A	output	A
	1		1		1
	1		1		1
	1		2		2
	2		2		
	3		4		
			5		

The number of copies of each tuple is the **minimum** of the number of copies in the subqueries

# SET OPERATORS: EXCEPT ALL

SELECT A FROM R  
**EXCEPT ALL**  
SELECT A FROM S;

R

A
1
1
1
2
3

S

A
1
1
2
2
4
5

output

A
1
3

The number of copies of each tuple is the **difference** (if positive) of the number of copies in the subqueries

# DISCUSSION ON DUPLICATES

---

- When doing projection:
  - easier to avoid eliminating duplicates
  - *tuple-at-a-time* processing
- When doing intersection, union or difference:
  - more efficient to **sort** the relations first
  - at that point you may as well eliminate the duplicates anyway

---

# NESTED QUERIES

---

# NESTED QUERIES

A parenthesized SELECT-FROM-WHERE statement (*subquery*) can be used as a value in a number of places:

- in **FROM** clauses
- in **WHERE** clauses

```
SELECT C.Name
FROM Country C
WHERE C.code =
      (SELECT C.CountryCode
       FROM City C
       WHERE C.name = 'Berlin');
```

Can you rewrite this query without a subquery (unnesting)?

# NESTING

---

- We can write nested queries because the SQL language is **compositional**
- Everything is represented as a multiset
- Hence the output of one query can be used as the input to another (**nesting**)



# NESTED QUERIES

---

*Find all countries in Europe with population more than 50 million*

```
SELECT C.Name
FROM (SELECT Name, Continent
      FROM Country
      WHERE Population >50000000) AS C
WHERE C.Continent = 'Europe' ;
```

Can you unnest this query?

---

# SET-COMPARISON OPERATOR: IN

---

*Find all countries in Europe that have **some** city with population more than 5 million*

```
SELECT C.Name
FROM Country C
WHERE C.Continent = 'Europe'
AND C.Code IN (SELECT CountryCode
                FROM City
                WHERE Population > 5000000);
```

# SET-COMPARISON OPERATOR: EXISTS


*Find all countries in Europe that have **some** city with population more than 5 million*

```
SELECT C.Name
FROM Country C
WHERE C.Continent = 'Europe'
AND EXISTS (SELECT *
             FROM City T
             WHERE T.Population > 5000000
             AND T.CountryCode = C.Code);
```

← correlated subquery

# SET-COMPARISON OPERATOR: ANY

*Find all countries in Europe that have **some** city with population more than 5 million*

```
SELECT C.Name
FROM Country C
WHERE C.Continent = 'Europe'
AND 5000000 <=  ANY (SELECT T.Population
      FROM City T
      WHERE T.CountryCode = C.Code);
```

# SET-COMPARISON OPERATORS

*Find all countries in Europe that have **all** cities with population less than 1 million*

```
SELECT C.Name
FROM Country C
WHERE C.Continent = 'Europe'
AND NOT EXISTS (SELECT *
                  FROM City T
                  WHERE T.Population > 1000000
                  AND T.CountryCode = C.Code);
```

# SET-COMPARISON OPERATORS: ALL

*Find all countries in Europe that have **all** cities with population less than 1 million*

```
SELECT C.Name
FROM Country C
WHERE C.Continent = 'Europe'
AND 1000000 > ALL (SELECT T.Population
                    FROM City T
                    WHERE T.CountryCode = C.Code);
```