

INDEXING

CS 564- Fall 2018

ACKs: Dan Suciu, Jignesh Patel, AnHai Doan

WHAT IS THIS LECTURE ABOUT?

- Indexes
 - alternative file organization
- Index classifications:
 - hash vs tree
 - clustered vs unclustered
 - primary vs secondary



FILE ORGANIZATION: RECAP

- So far we have seen **heap files**
 - they store **unordered data**
 - fast for scanning all records in a file
 - fast for retrieving by **record id (rid)**
- But we also need alternative organizations of a file to support **other access patterns**

MOTIVATION

- Consider the following SQL query:

SELECT *

FROM Sales

WHERE Sales.date = “02-11-2016”

- For a heap file, we have to scan all the pages of the file to return the correct result

ALTERNATIVE FILE ORGANIZATIONS

- We can *speed up* the query execution by better organizing the data in a file
- There are many alternatives:
 - sorted files
 - indexes
 - B+ tree
 - hash index
 - bitmap index

INDEX BASICS

WHAT IS AN INDEX?

- **Index**: a data structure that organizes records of a table to optimize retrieval
 - it speeds up searches for a subset of records, based on values in certain (*search key*) attributes
 - any subset of the fields of a relation can be the search key
 - a search key is *not* the same as the primary key!
- An index contains a collection of *data entries* (each entry with enough info to locate the records)

HASH INDEX: EXAMPLE

- A hash index is a collection of *buckets*
 - bucket = primary page + overflow pages
 - each bucket contains one or more data entries
- To find the bucket for each record, we use a hash function h applied on the search key k
 - N = number of buckets
 - $h(k) \bmod N$ = bucket in which the data entry belongs
- Records with different search key may belong in the same bucket!

HASH INDEX: EXAMPLE

Person(name, zipcode, phone)

- *search key*: zipcode
- *hash function h* : last 2 digits 

- 4 buckets

- each bucket has 2 data entries (full record)

overflow pages

bucket 0

(John, 534**00**, 23218564)
(Alice, 547**68**, 60743111)



(Anna, 536**32**, 23209964)

bucket 1

(Paris, 534**09**, 23200564)


bucket 2



bucket 3


(Maria, 344**11**, 29010533)

data entry



DATA ENTRIES

The actual data may not be in the same file as the index!


- In a data entry with search key **k** we have three alternatives of what to store:
 1. the record with key value **k**
 2. **<k, rid of record with search key value k>**
 3. **<k, list of rids of records with search key k>**
- The choice of alternative for data entries is **independent** of the indexing technique 

EXAMPLE

Person(name, zipcode, phone)


- *search key*: zipcode
- *hash function h* : last 2 digits

data entry



ALTERNATIVES FOR DATA ENTRIES

Alternative #1: *the data entry contains the record*

- the index structure is *by itself* a file organization for records
- *at most one* index on a given collection of data records should use alternative #1 
- if data records are very large, the number of pages containing data entries is high
 - this means possibly *slower search!*

ALTERNATIVES FOR DATA ENTRIES

Alternatives #2, #3: *the data entry contains the rid*

- Data entries are typically much smaller than data records. So, better than #1 with large data records, especially if search keys are small
- #3 is more compact than #2, but leads to variable sized data entries even if search keys are of fixed length

MORE ON INDEXES

A file can have several indexes, on different search keys!

Index classification:

- *primary* **vs** *secondary*
- *clustered* **vs** *unclustered*

PRIMARY VS SECONDARY

- If the search key contains the **primary key**, it is called a **primary index**
 - in a primary index, there are no duplicates for a value of the search key
 - there can only be one primary index!
- Any other index is called a **secondary index**
- If the search key contains a **candidate key**, it is called a **unique index**
 - a unique index can also return no duplicates


EXAMPLE

Sales (sid, product, date, price)

1. An index on (sid) is a primary and unique index
2. An index on (date) is a secondary, but not unique, index

CLUSTERED INDEXES

Clustered index: the order of records **matches** the order of data entries in the index

- alternative #1 implies that the index is clustered
- a table can have at most one clustered index 
- the cost of retrieving data records through the index varies greatly based on whether index is clustered or not

logical order of index ~ physical order of records

INDEXES IN PRACTICE

CHOOSING INDEXES

- What indexes should we create?
 - which relations should have indexes?
 - what field(s) should be the search key?
 - should we build several or one index?
- For each index, what kind of an index should it be?
 - clustered
 - hash/tree/bitmap

CHOOSING INDEXES

- Consider the best plan using the current indexes, and see if a better plan is possible with an additional index
- One must understand how a DBMS evaluates queries and creates query evaluation plans
- Important trade-offs:
 - queries go faster, updates are slower
 - more disk space is required

CHOOSING INDEXES

- Attributes in **WHERE** clause are candidates for index keys
 - exact match condition suggests hash index
 - indexes also speed up joins (later in class)
 - range query suggests tree index (B+ tree)
- Multi-attribute search keys should be considered when a **WHERE** clause contains several conditions
 - order of attributes is important for range queries
 - such indexes can enable **index-only** strategies for queries

COMPOSITE INDEXES

Composite search keys: search on a combination of fields (e.g. <date, price>)

- **equality query:** every field value is equal to a constant value
 - date="02-20-2015" and price =75
- **range query:** some field value is not a constant
 - date="02-20-2015"
 - date="02-20-2015" and price > 40

INDEXES IN SQL

```
CREATE INDEX index_name  
ON table_name (column_name);
```

- Example of simple search key:

```
CREATE INDEX index1  
ON Sales (price);
```

INDEXES IN SQL

```
CREATE UNIQUE INDEX index2  
ON Sales (sid);
```

- A unique index does not allow any duplicate values to be inserted into the table
- It can be used to check efficiently integrity constraints (a duplicate value will not be allowed to be inserted)

INDEXES IN SQL

```
CREATE INDEX index3  
ON Sales (date, price);
```

- Indexes with composite search keys are larger and more expensive to update
- They can be used if we have multiple selection conditions in our queries