

# THE B+ TREE INDEX

---

*CS 564- Fall 2018*

---

*ACKs: Jignesh Patel, AnHai Doan*

---

# WHAT IS THIS LECTURE ABOUT?

---

## The **B+ tree** index

- Basics
- Search/Insertion/Deletion
- Design & Cost

# INDEX RECAP

---

- We have the following query:

```
SELECT  *  
FROM    Sales  
WHERE   price > 100 ;
```

- How do we organize the file to answer this query efficiently?

# INDEXES

---

- Hash index:
  - good for equality search
  - in expectation constant I/O cost for search and insert
- B+ tree index:
  - good for **range** and **equality** search

---

# B+ TREE BASICS


---

# THE B+ TREE INDEX

---

- a dynamic tree-structured index
  - adjusted to be always height-balanced
  - 1 node = 1 physical page
- supports efficient equality and range search
- widely used in many DBMSs
  - SQLite uses it as the default index
  - SQL Server, DB2, ...

# B+ TREE INDEX: BASIC STRUCTURE



## data entries

- exist *only* in the leaf nodes
- are sorted according to the search key

# B+ TREE: NODE




- parameter  $d$  is the *order* of the tree
- each node contains  $d \leq m \leq 2d$  entries
  - minimum 50% occupancy at all times
- with the exception of the **root node**, which can have  $1 \leq m \leq 2d$  entries





# NON-LEAF NODES

An non-leaf (or internal) node with  $m$  entries has  $m+1$  pointers to lower-level nodes




# LEAF NODES

A leaf node with  $m$  entries has



- $m$  pointers to the data records (rids)
- pointers to the **next** and **previous** leaves



---

# B+ TREE OPERATIONS

---

---

# B+ TREE OPERATIONS

---

A B+ tree supports the following operations:

- equality search
- range search
- insert
- delete
- bulk loading

# SEARCH

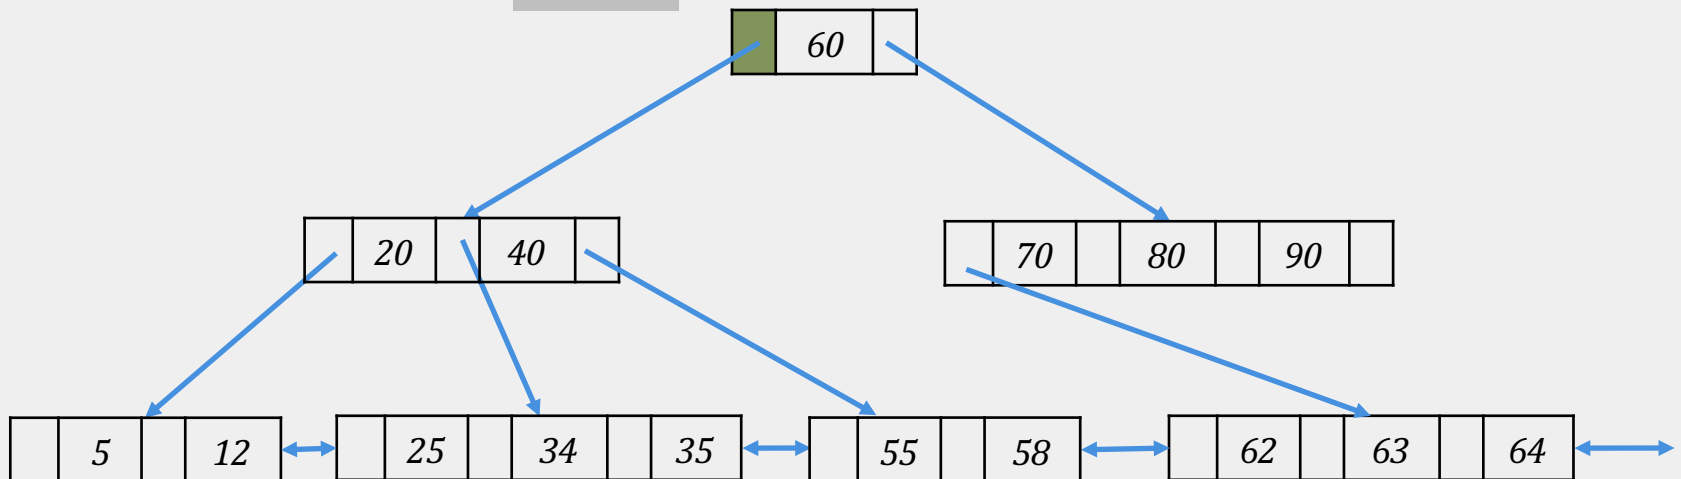
---

- start from the root node
- examine the index entries in non-leaf nodes to find the correct child
- traverse down the tree until a leaf node is reached
  - for equality search, we are done
  - for range search, traverse the leaves sequentially using the previous/next pointers

# EQUALITY SEARCH: EXAMPLE

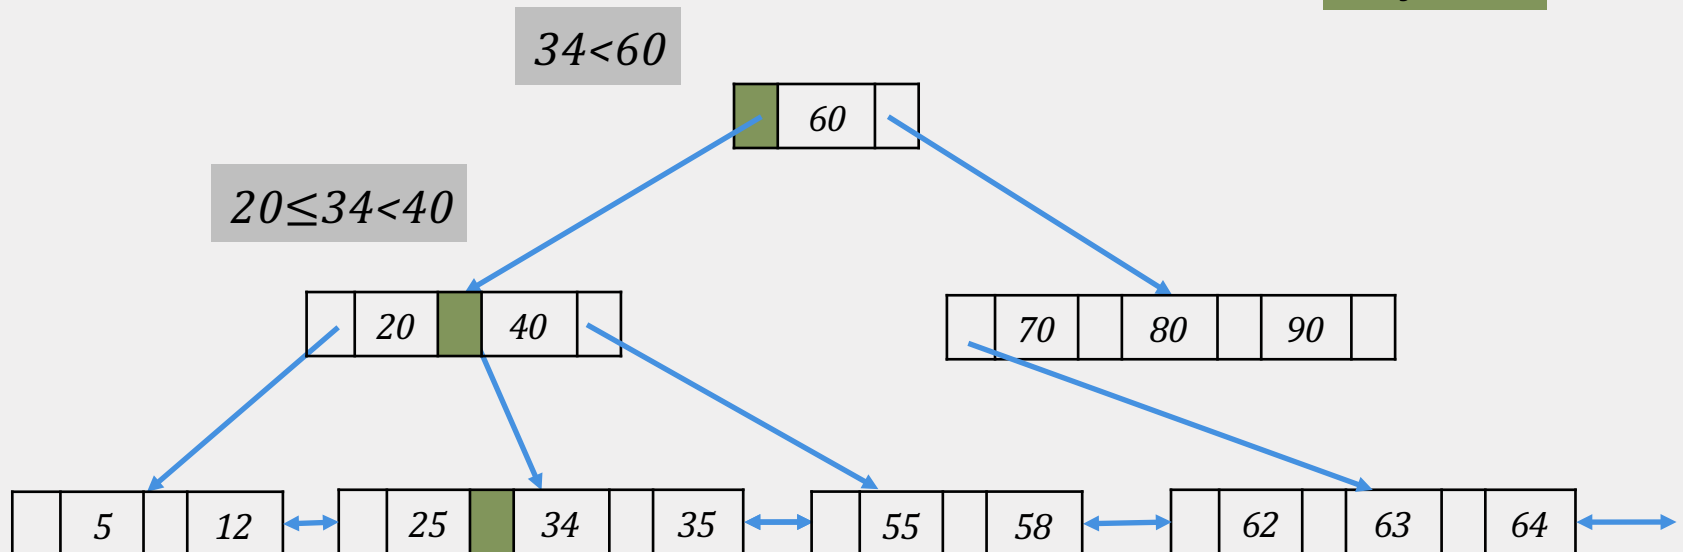
key = 34

$34 < 60$



# EQUALITY SEARCH: EXAMPLE

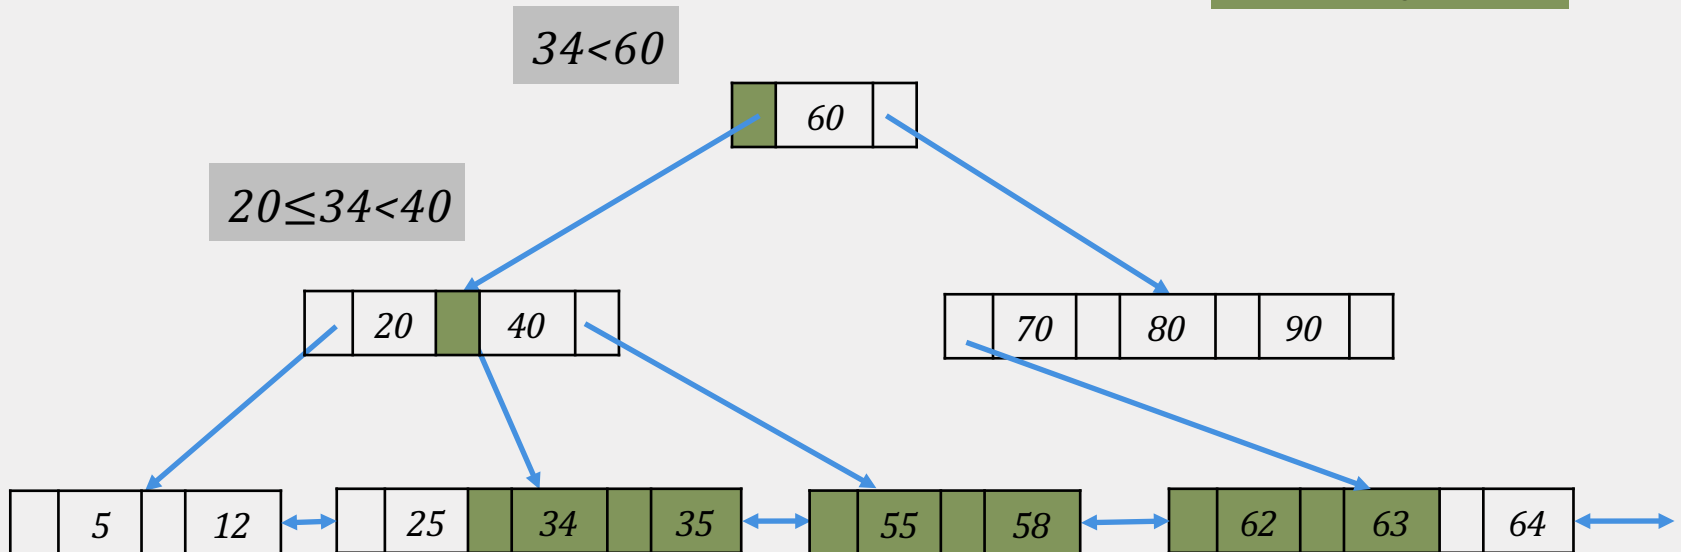
key = 34



To locate the correct data entry in the leaf node, we can do either linear or binary search

# RANGE SEARCH: EXAMPLE

$34 \leq \text{key} \leq 63$



After we find the leftmost point of the range,  
we traverse sequentially!




# INSERT

- find correct leaf node **L**
- insert data entry in **L**
  - If **L** has enough space, DONE!
  - Else, we must **split** **L** (into **L** and a new node **L'**)
    - redistribute entries evenly, **copy up** the middle key
    - insert index entry pointing to **L'** into parent of **L**
- This can propagate **recursively** to other nodes!
  - to split a non-leaf node, redistribute entries evenly, but **push up** the middle key

# INSERT: EXAMPLE

order  $d = 2$

insert 8





# INSERT: EXAMPLE

order  $d = 2$

insert 8

the leaf node is full so  
we must split it!



d entries


d+1 entries

# INSERT: EXAMPLE

order  $d = 2$

insert 8


the middle key (5) must be copied up,  
but the root node is full as well!



# INSERT: EXAMPLE

order  $d = 2$

insert 8



# INSERT PROPERTIES

---

The B+ Tree insertion algorithm has several attractive qualities:

- ~ same cost as exact search
- it is ***self-balancing***: the tree remains balanced (with respect to height) even after multiple insertions


# B+ TREE: DELETE

- find leaf node **L** where entry belongs
- remove the entry
  - If **L** is at least half-full, DONE!
  - If **L** has only  $d-1$  entries,
    - Try to **re-distribute**, borrowing from **sibling**
    - If re-distribution fails, **merge** **L** and sibling
- If a **merge** occurred, we must **delete** an entry from the parent of **L**

# DELETE : EXAMPLE 1

order  $d = 2$

delete 22




since by deleting 22 the node remains half-full, we simply remove it



# DELETE : EXAMPLE 1

order  $d = 2$


delete 22



# DELETE : EXAMPLE 2

order  $d = 2$

delete 20

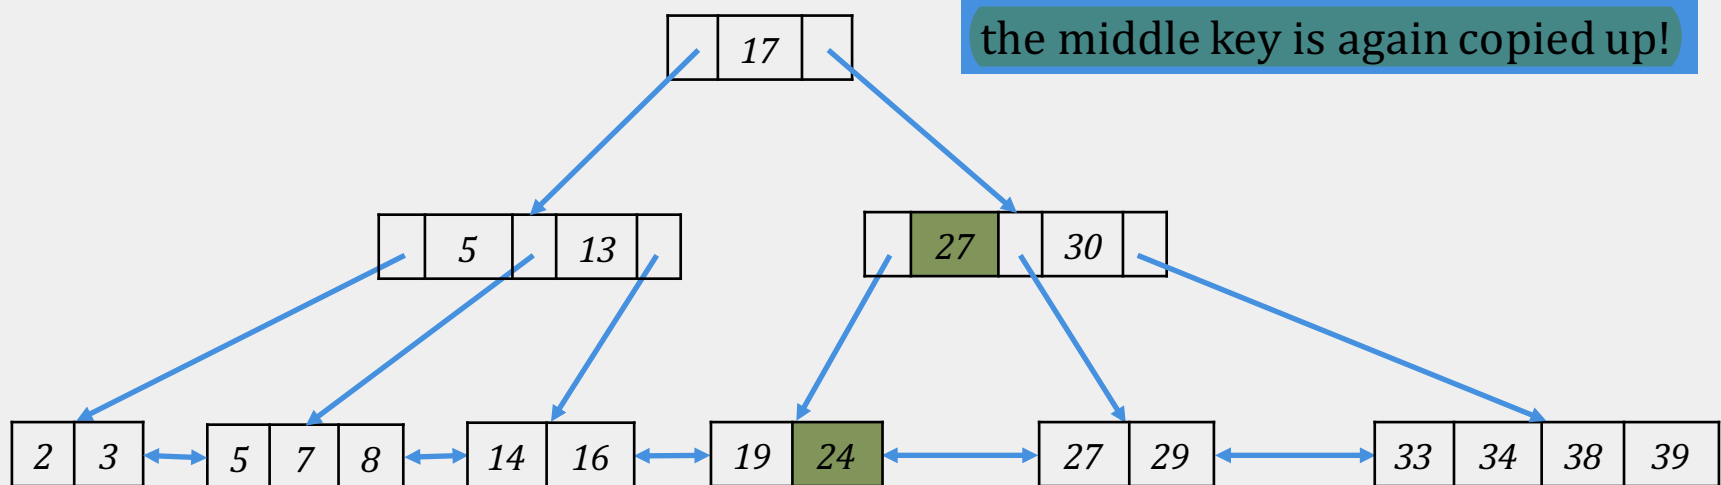


by removing 20 the node is not half-full anymore,  
so we attempt to redistribute!

# DELETE : EXAMPLE 2

order  $d = 2$

delete 20

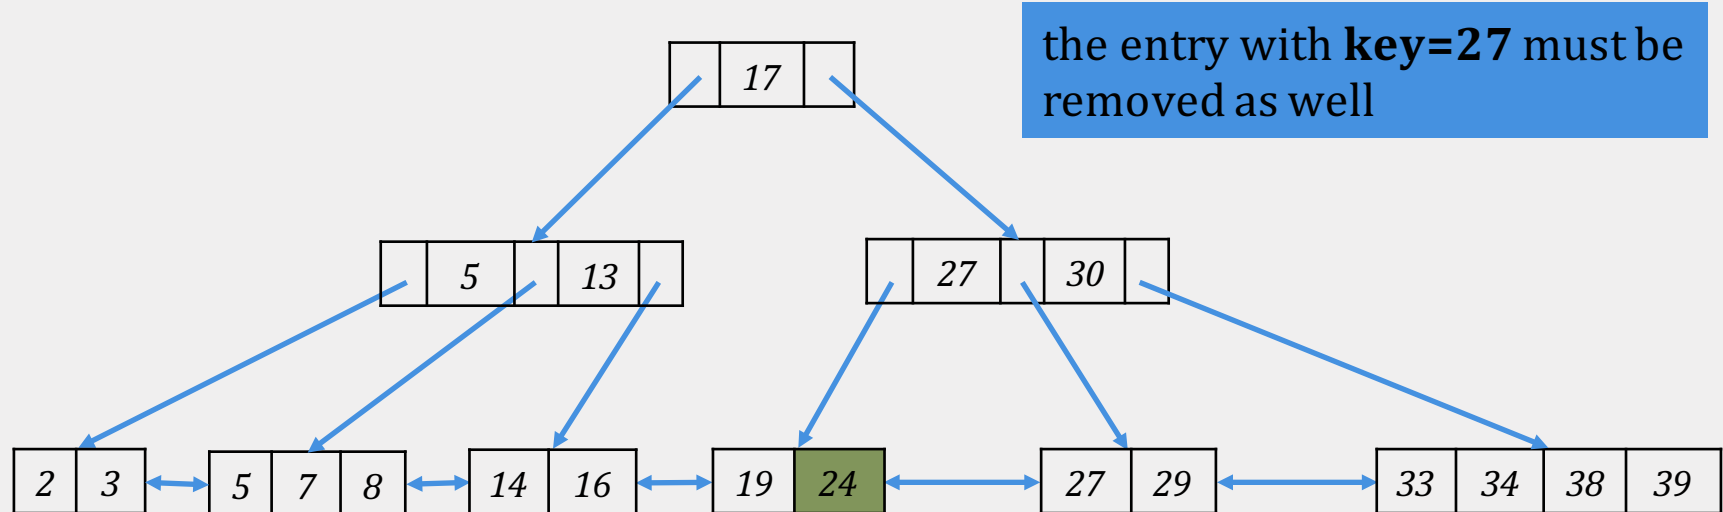


by removing 20 the node is not half-full anymore, so we attempt to redistribute!

# DELETE : EXAMPLE 3

order  $d = 2$

delete 24




in this case, we have to merge nodes!

# DELETE : EXAMPLE 3

order  $d = 2$

delete 24




# DELETE : EXAMPLE 3

order  $d = 2$

delete 24

we are not done, since the resulting non-leaf node is not half-full!



# B+ TREE: DELETE

---

- Redistribution of entries can also be possible for the non-leaf nodes
- We can also try to redistribute using *all siblings*, and not only the neighboring one

# DUPLICATES

---

- **duplicate keys**: many data entries with the same key value
- Solution 1:
  - All entries with a given key value reside on a single page
  - Use overflow pages
- Solution 2:
  - Allow duplicate key values in data entries
  - Modify search operation



---

# **B+ TREE DESIGN & COST**

---

# B+ TREE: FAN-OUT

**fan-out**  $f$ : the number of pointers to child nodes coming out of a non-leaf node

- compared to binary trees (fan-out = 2), B+ trees have a high fan-out ( $d+1 \leq f \leq 2d+1$ )
- The fan-out of B+ trees is dynamic, but we will often assume it is constant for our cost model

# B+ TREE: FILL-FACTOR

---

**fill-factor**  $F$ : the percent of available slots in the B+ Tree that are filled

- it is usually  $< 1$  to leave slack for (quicker) insertions!
- typical fill factor  $F = 2/3$

# B+ TREE: HEIGHT

**height  $h$ :** the number of levels of the non-leaf nodes

- the height is at least 1 (root node)
- high fan-out  $\rightarrow$  smaller height  $\rightarrow$  less I/O per search
- typical heights of B+ trees: 3 or 4

# B+ TREE: EXAMPLE

---

- page size  $P = 4000$  bytes
- search key size = 30 bytes
- address size = 10 bytes
- fill-factor  $F = 2/3$
- number of records = 2,000,000
  
- We assume that the data entries store only the search key and the address of tuple
- We assume no duplicate entries

# B+ TREE: EXAMPLE

What is the order  $d$  and fan-out  $f$ ?

- each non-leaf node stores up to  $2d$  values of the key +  $(2d+1)$  addresses for the children pages
- to fit this into a single page, we must have:

$$2d \cdot 30 + (2d + 1) \cdot 10 \leq 4000$$

$$d \leq 50$$

- since a maximum capacity node has  $(2d+1) = 101$  children, and the fill-factor is  $2/3$ , the fan-out is  $f = 101 * \frac{2}{3} = 67$

# B+ TREE: EXAMPLE

---

How many leaf pages are in the B+ tree?

- we assume for simplicity that each leaf page stores only pairs of (key, address)
- each pair needs  $30+10 = 40$  bytes
- to store 2,000,000 such pairs with fill-factor  $F = 2/3$ , we need:

$$\#leaves = (2,000,000 * 40) / (4,000 * F) = 30,000$$

# B+ TREE: EXAMPLE

What is the height  $h$  of the B+ tree?

- we calculated that we need to index  $N = 30,000$  pages
- $h = 1$  -> indexes  $f$  pages
- $h = 2$  -> indexes  $f^2$  pages
- ...
- $h = k$  -> indexes  $f^k$  pages

height must be  $h = \lceil \log_f N \rceil$

for our example,  $h = \lceil \log_{67} 30,000 \rceil = 3$



# B+ TREE: EXAMPLE

---

What is the total size of the tree?

- $\#pages = 1 + 67 + 67^2 + 30,000 = 34,557$
- the top levels of the B+ tree do not take much space and can be kept in the buffer pool
  - *level 0* = 1 page ~ 4 KB
  - *level 1* = 67 pages ~ 268 KB
  - *level 2* = 4,489 pages ~ 18 MB

# COST MODEL FOR SEARCH

To do equality search:

- we read one page per level of the tree
- levels that we can fit in buffer are free!
- finally we read in the actual record



$I = 0$  if the record is stored at the leaf node, otherwise  $I = 1$

$$\text{I/O cost} = h - L_B + 1 + I$$

If we have  $B$  available buffer pages, we can store  $L_B$  levels of the B+ Tree in memory:

- $L_B$  is the number of levels such that the sum of all the levels' nodes fit in the buffer:

$$B \geq 1 + f + \dots + f^{L_B - 1}$$

# COST MODEL FOR SEARCH

To do range search:

- we read one page per level of the tree
- levels that we can fit in buffer are free!
- we read sequentially the pages in the range

$$\text{I/O cost} = h - L_B + OUT$$

Here,  $OUT$  is the I/O cost of loading the additional leaf nodes we need to access + the I/O cost of loading each *page* of the results