

# EXTERNAL SORTING

---

*CS 564- Fall 2018*

---

*ACKs: Dan Suciu, Jignesh Patel, AnHai Doan*

---

# WHAT IS THIS LECTURE ABOUT?

---

I/O aware algorithms for **sorting**

- External merge
  - a primitive for sorting
- External merge-sort
  - basic algorithm
  - optimizations

# WHY SORTING?

---

- users often want the data sorted (**ORDER BY**)
- first step in bulk-loading a B+ tree
- used in duplicate elimination (why?)
- the **sort-merge join** algorithm (later in class) involves sorting as a first step

# **SORTING IN DATABASES**

Why don't the standard sorting algorithms work for a database system?

- merge sort
- quick sort
- heap sort

**The data typically does not fit in memory!**

e.g. how do we sort 1TB of data with 8GB of RAM?

---

# EXTERNAL MERGE

---

# EXTERNAL MERGE PROBLEM


**Input:** 2 sorted lists (with  $M$  and  $N$  pages)

**Output:** 1 **merged** sorted list (with  $M+N$  pages)

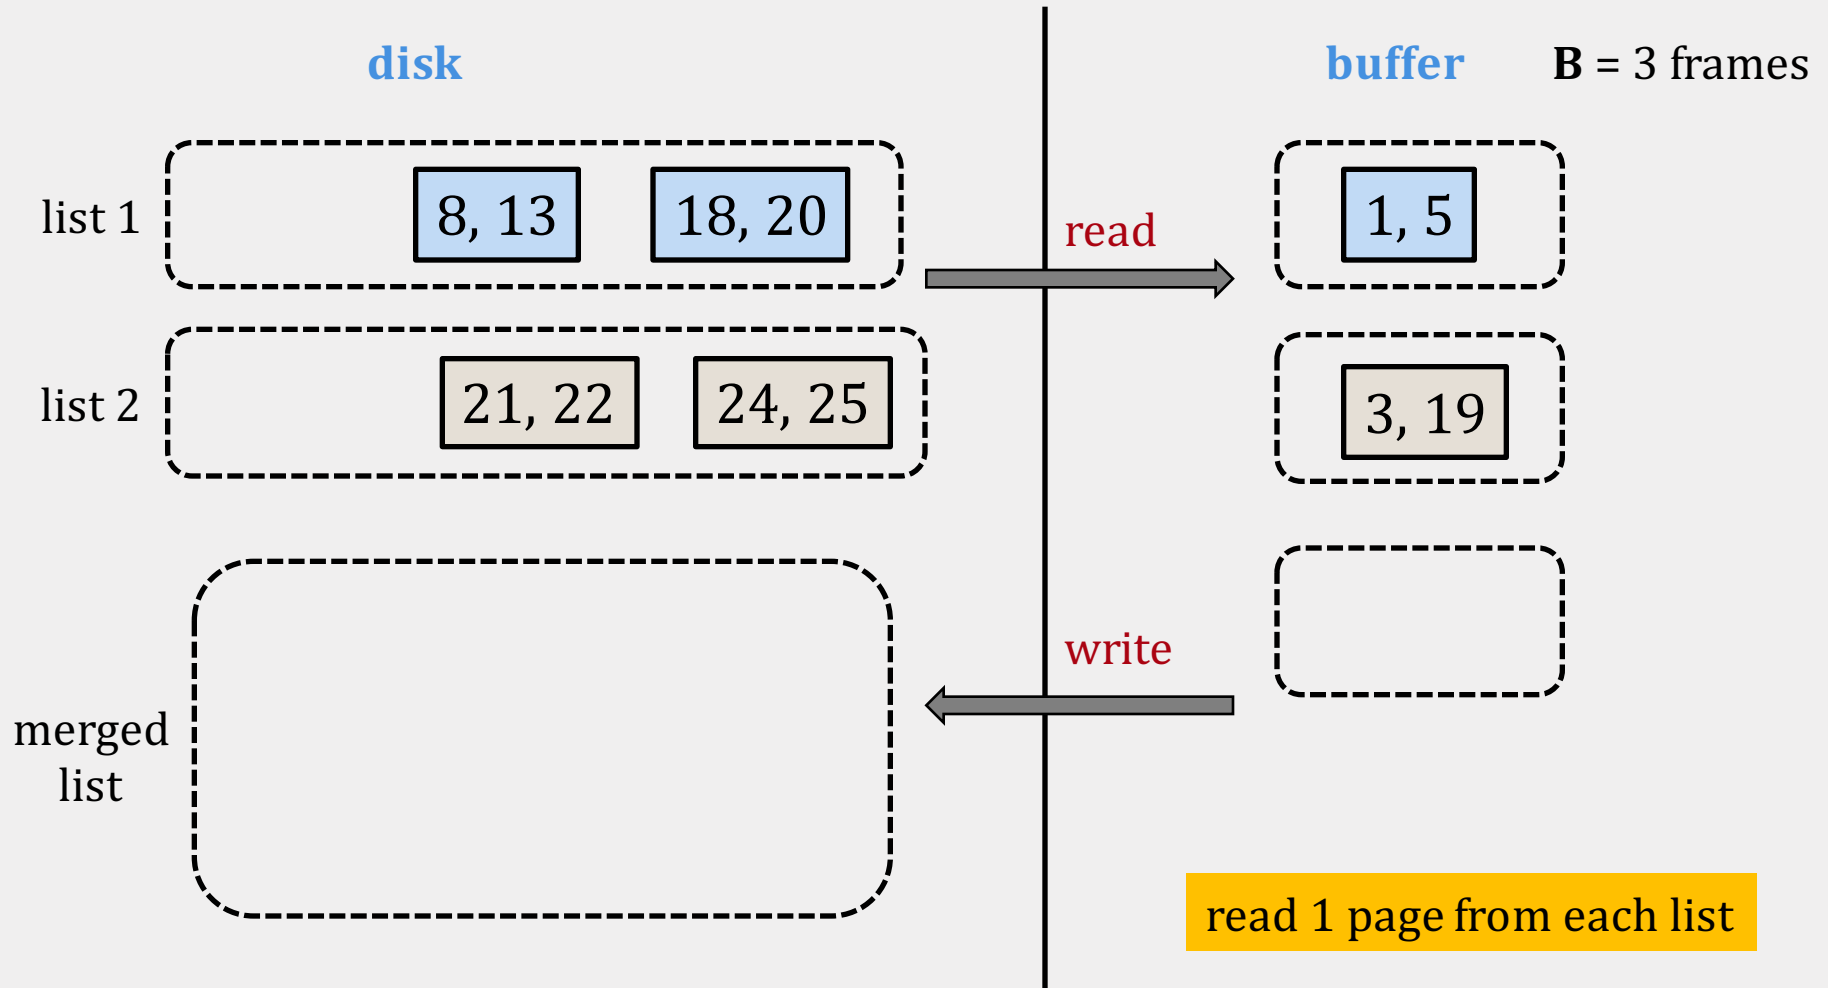
Can we efficiently (in terms of I/O) merge the two lists using a buffer of size at least 3?

Yes, using only  $2(M+N)$  I/Os !

# EXTERNAL MERGE ALGORITHM

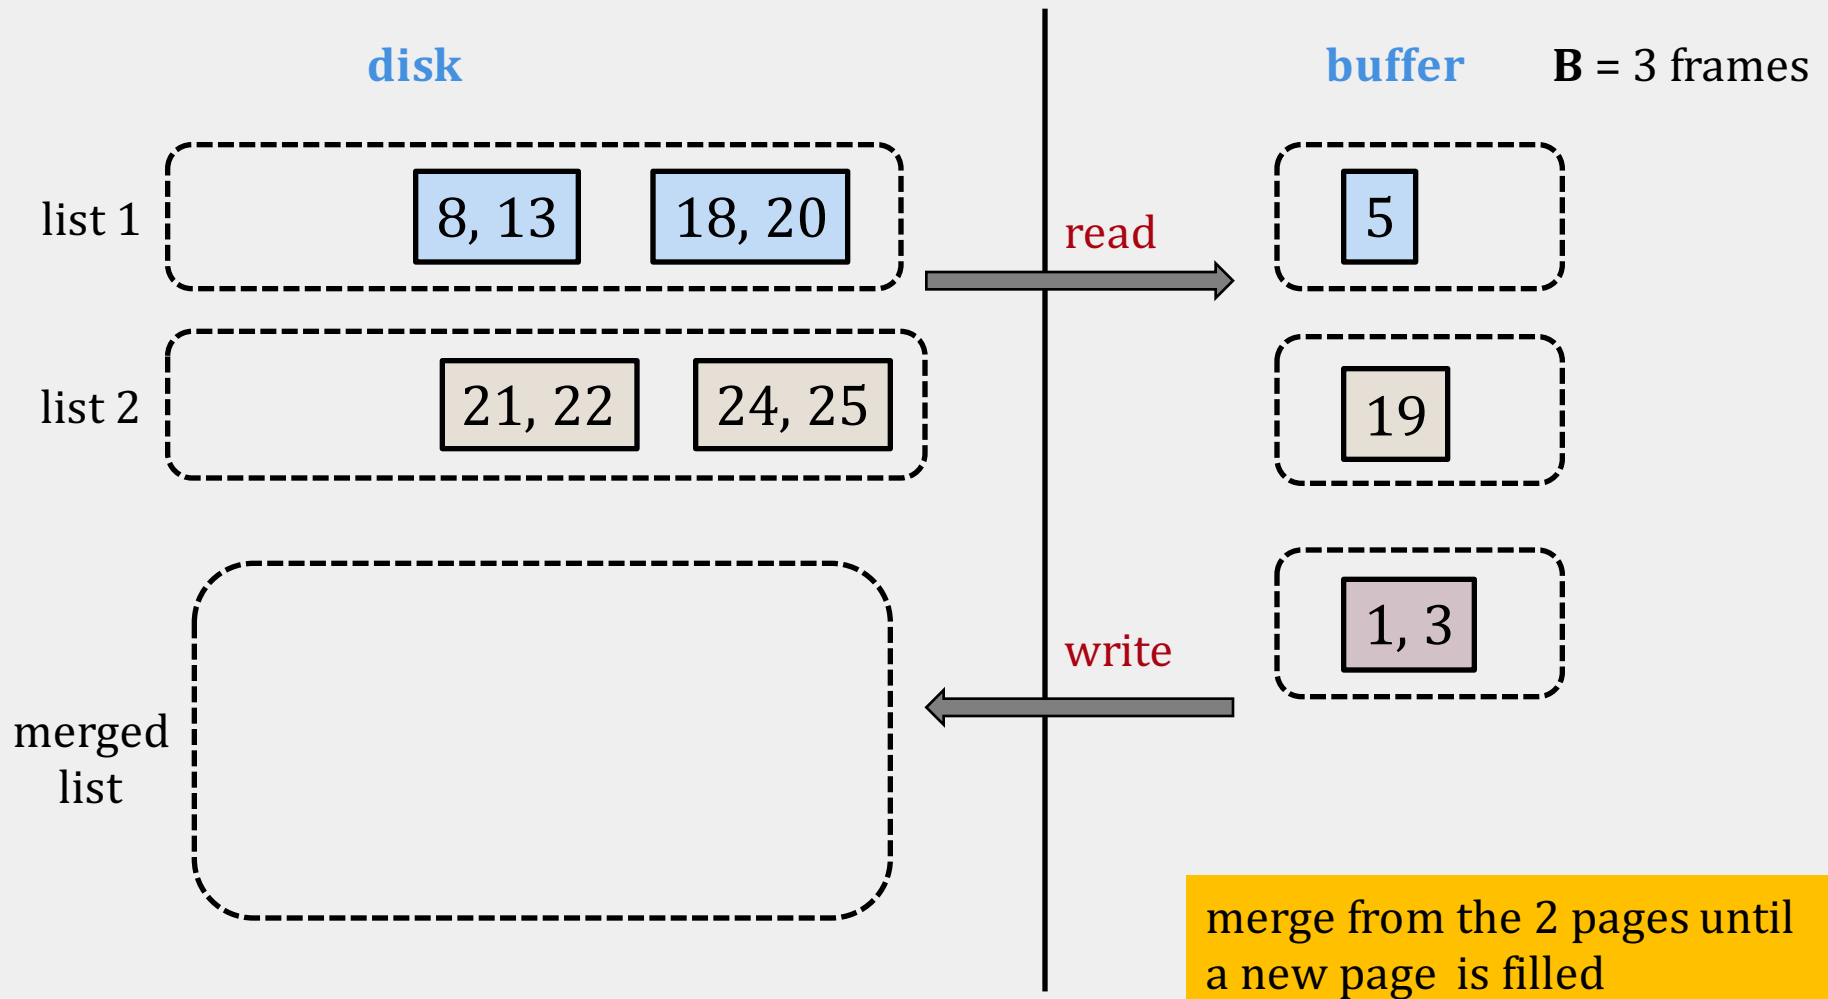


# EXTERNAL MERGE ALGORITHM

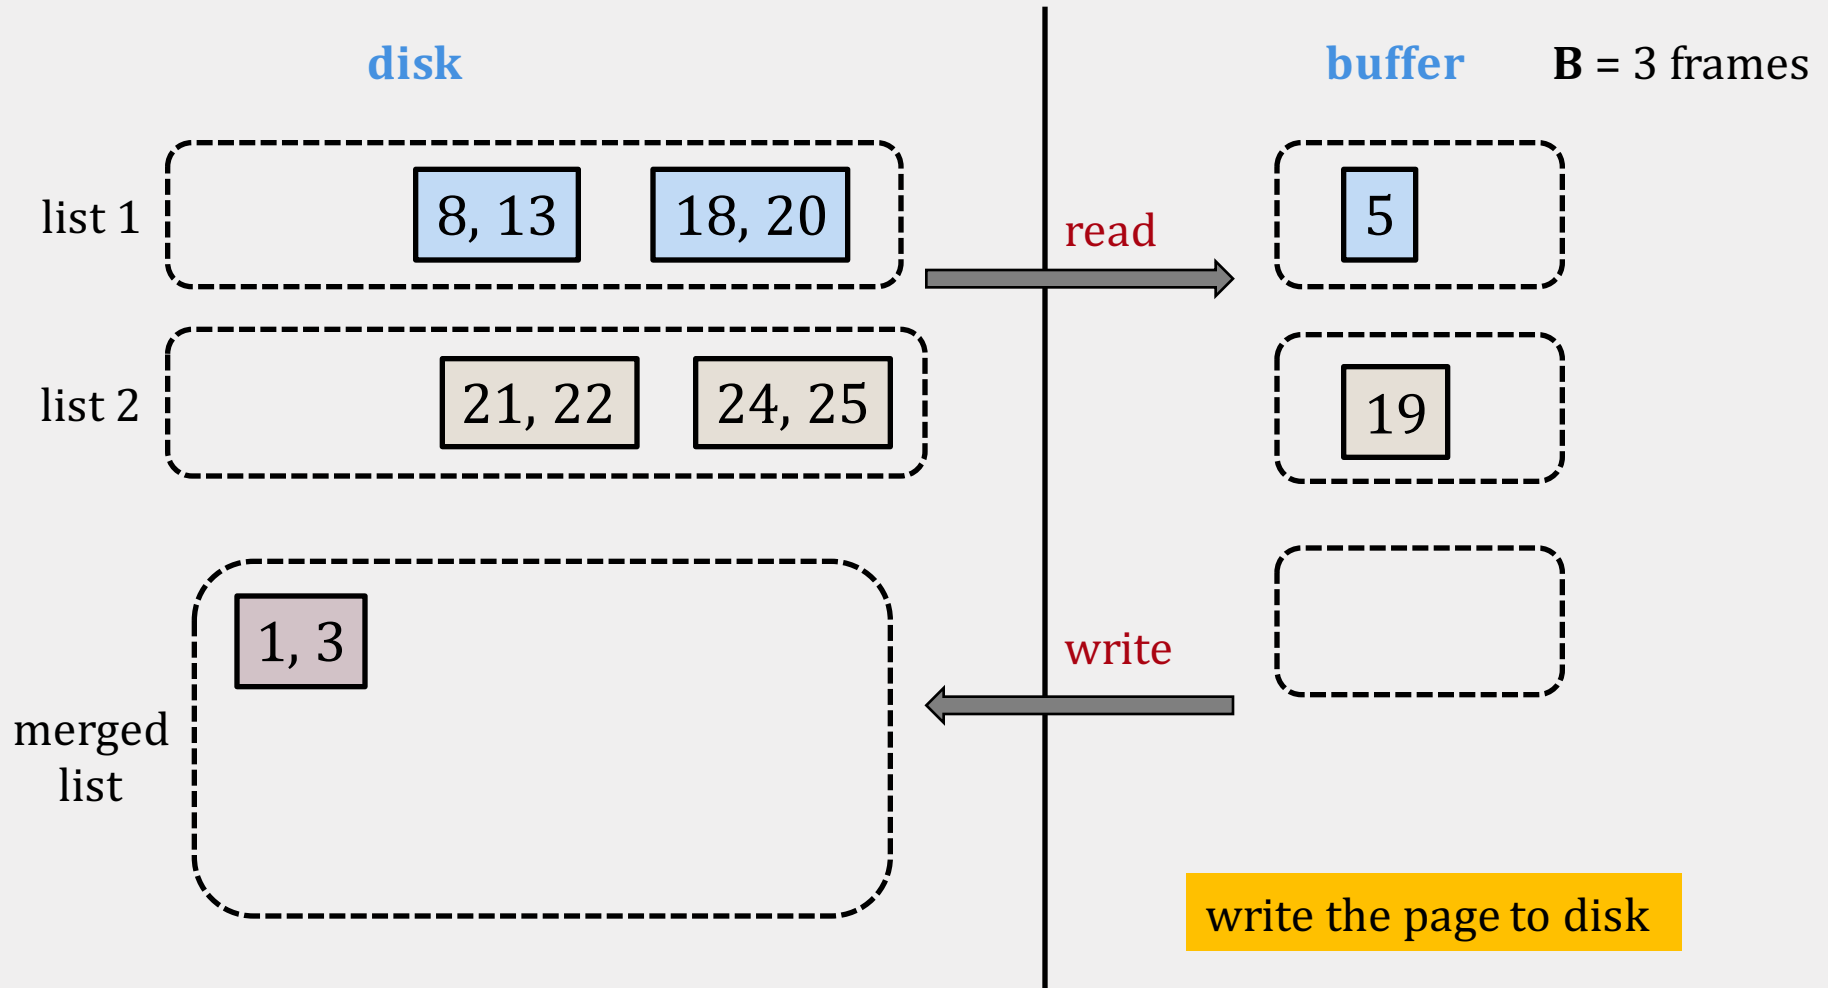




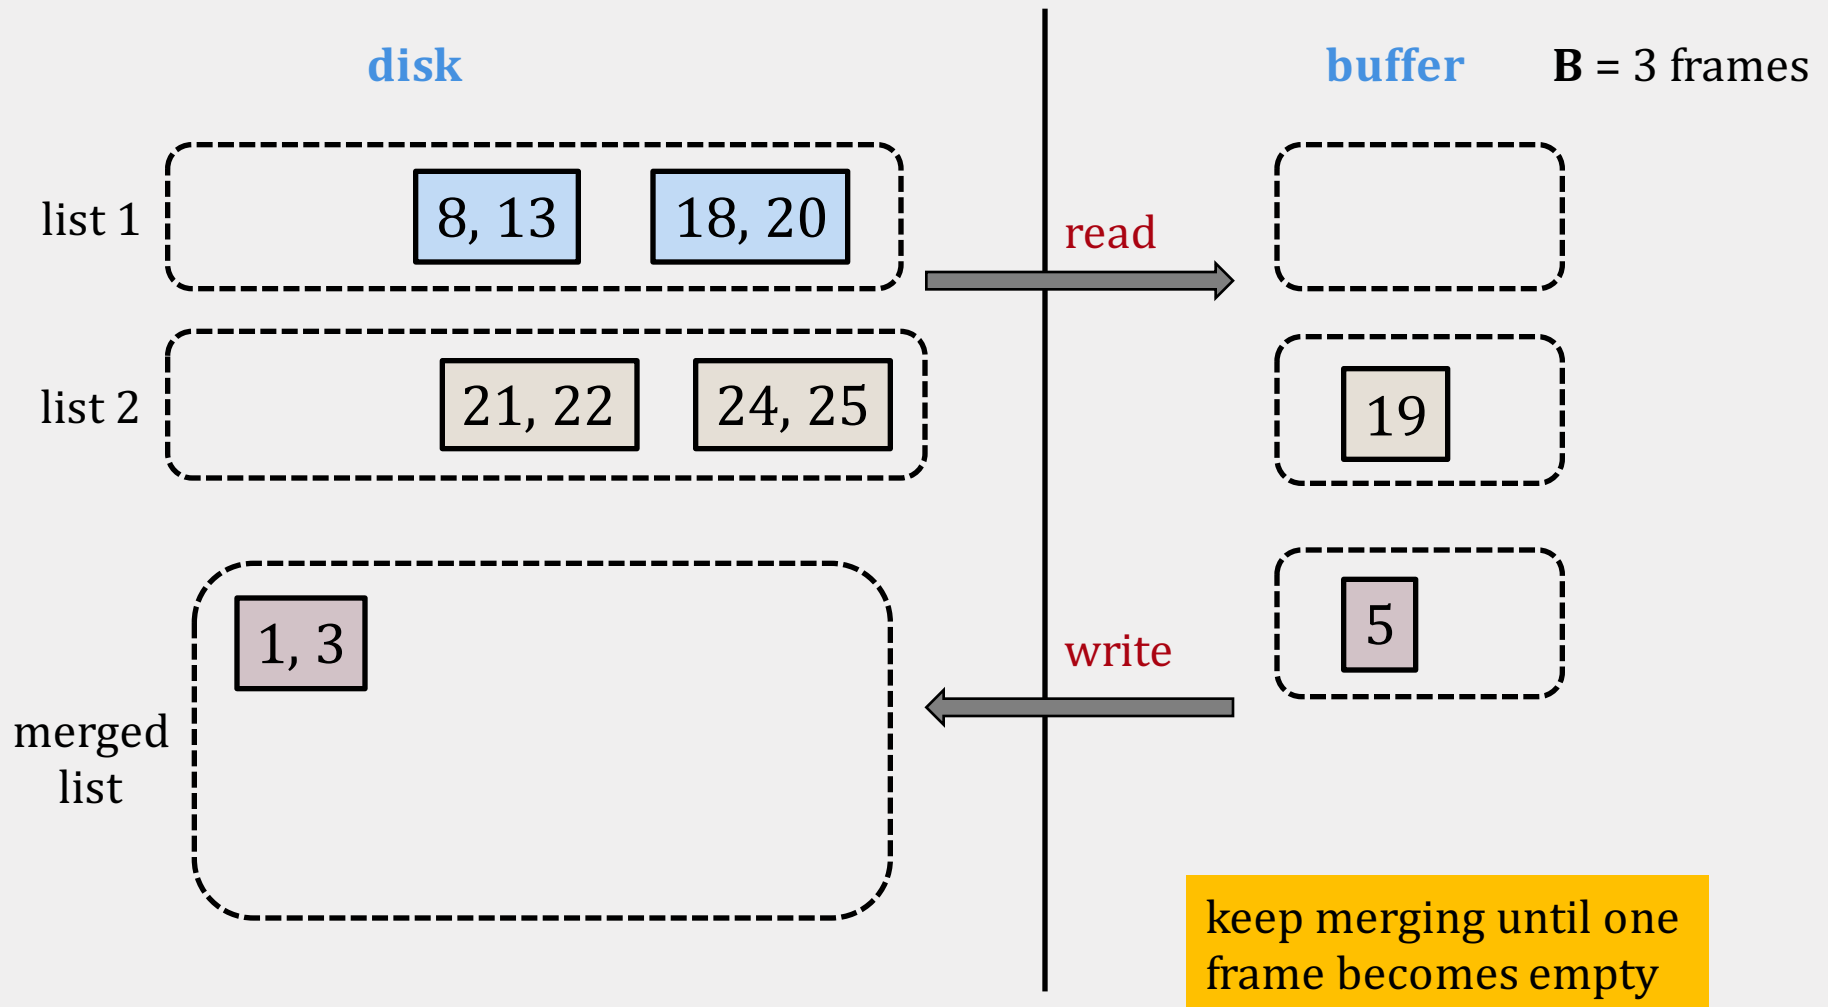
# EXTERNAL MERGE ALGORITHM



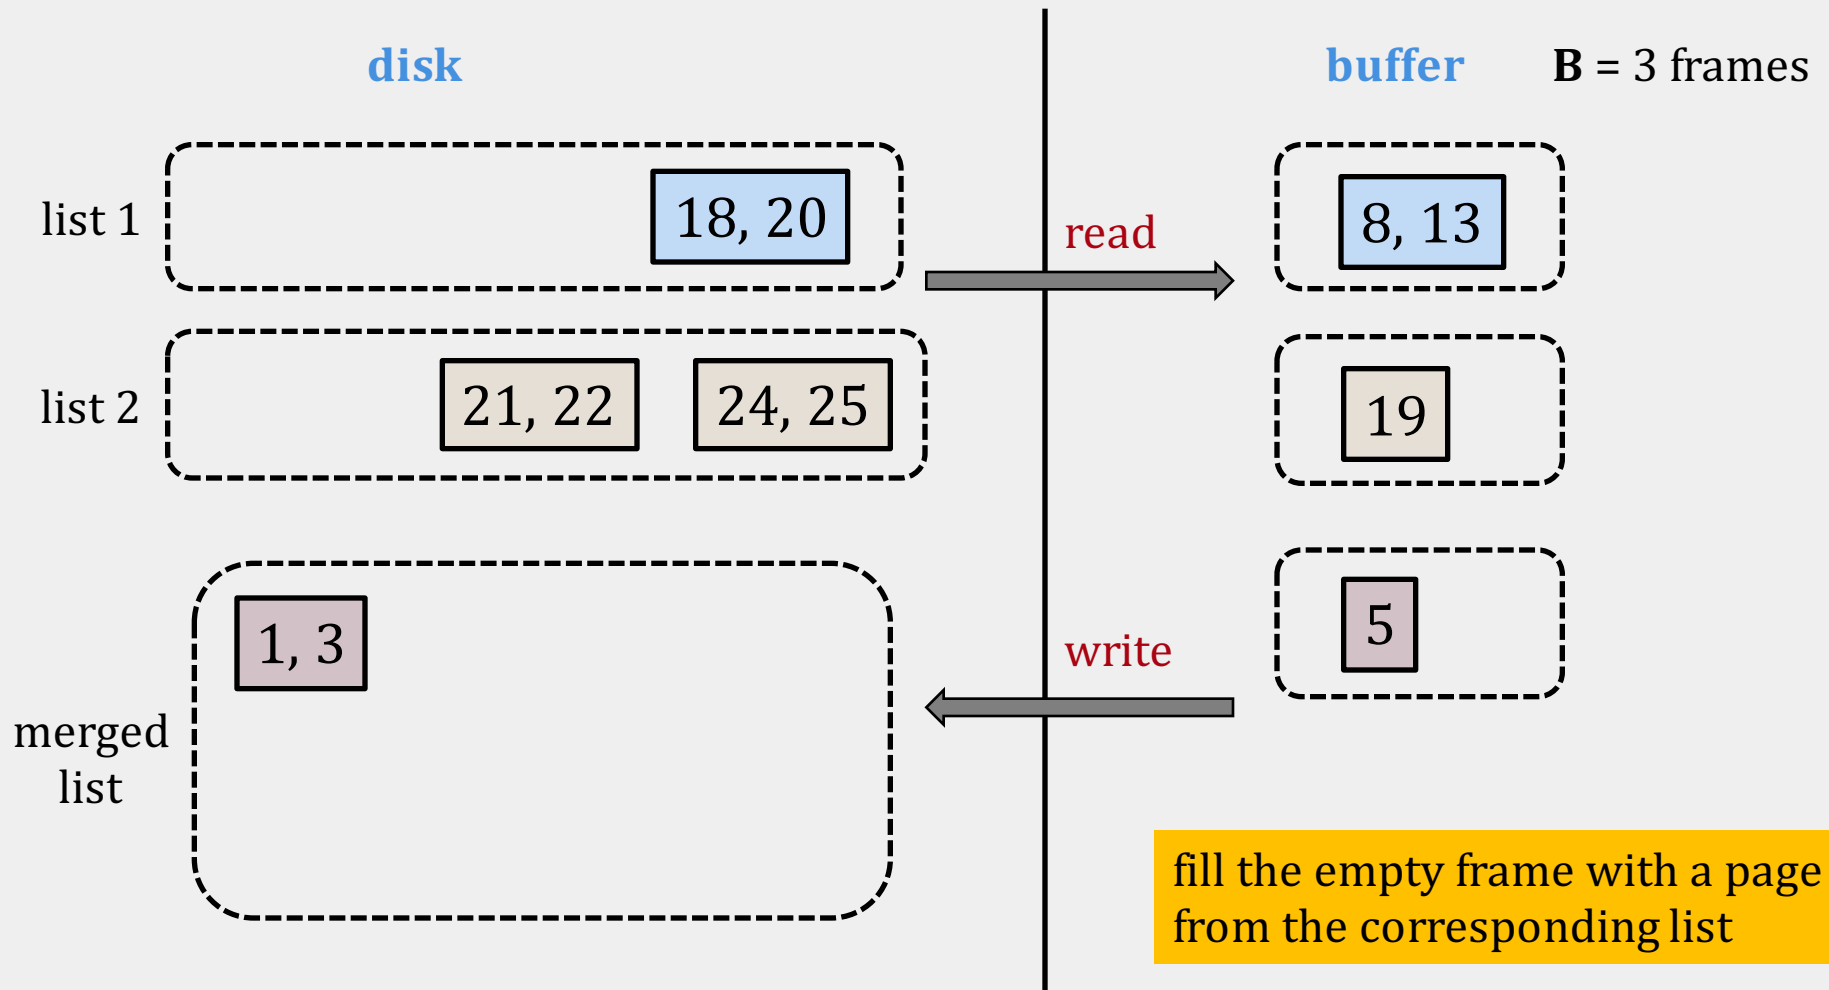
# EXTERNAL MERGE ALGORITHM



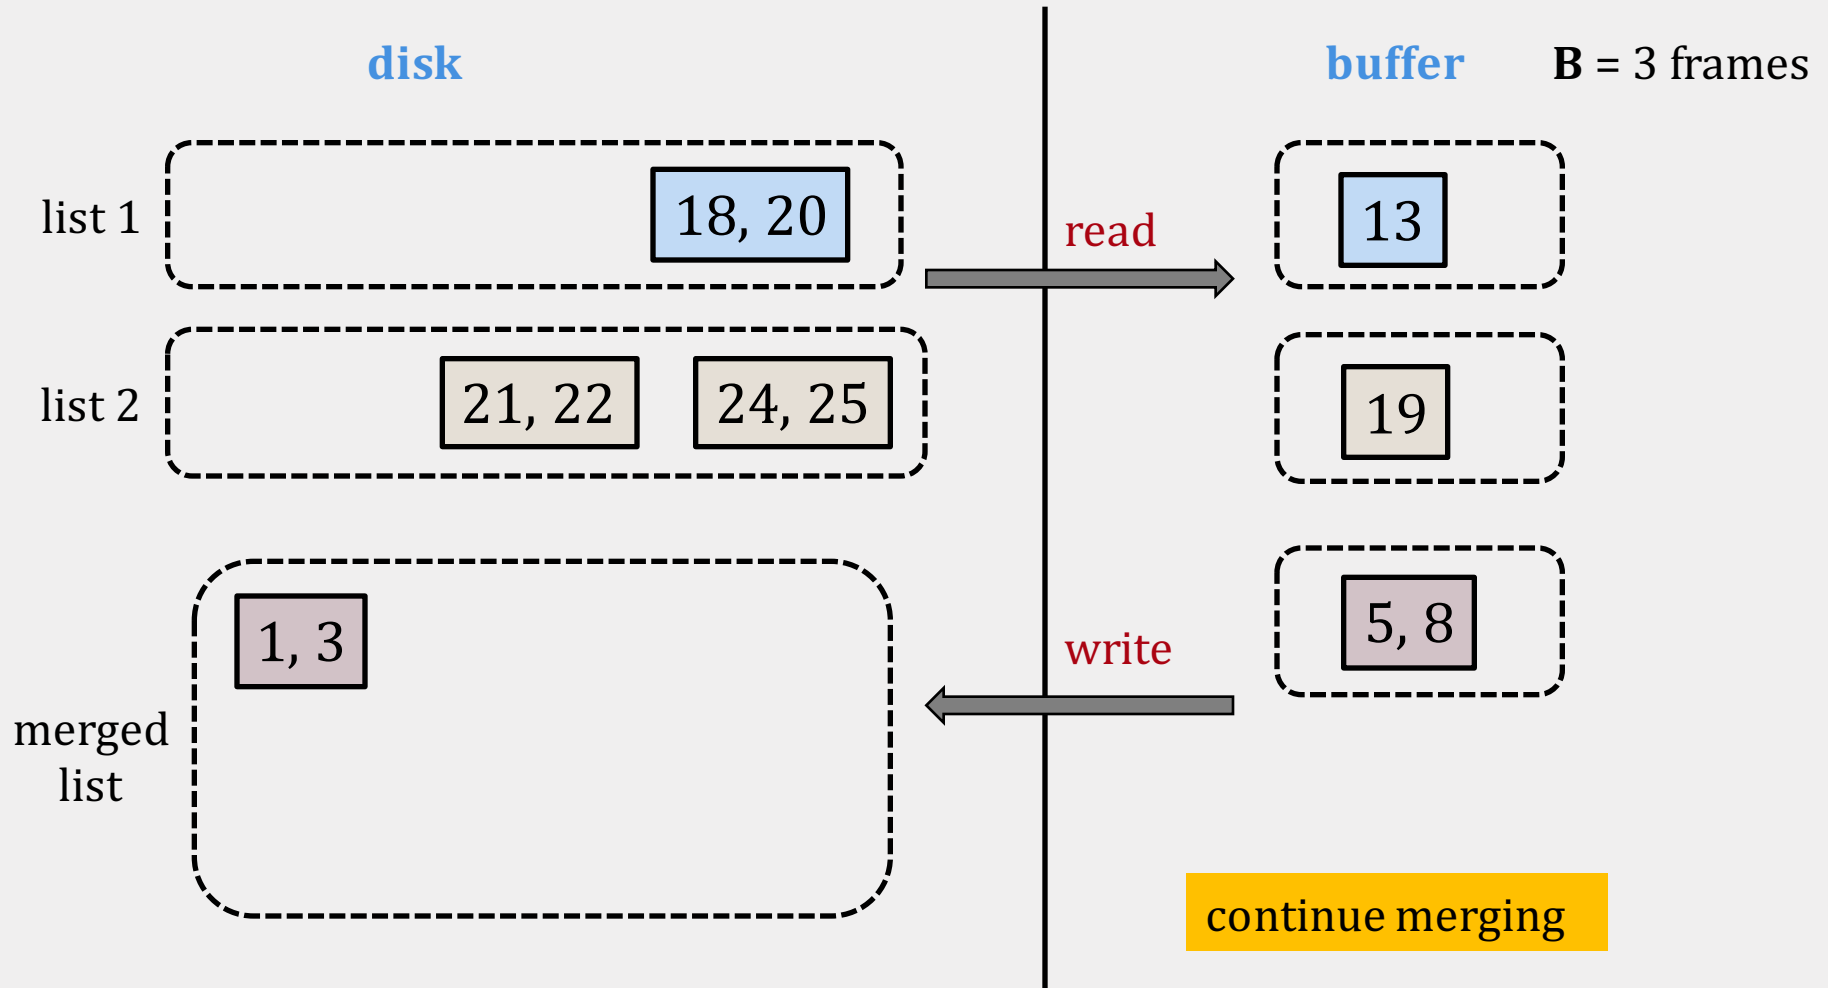
# EXTERNAL MERGE ALGORITHM



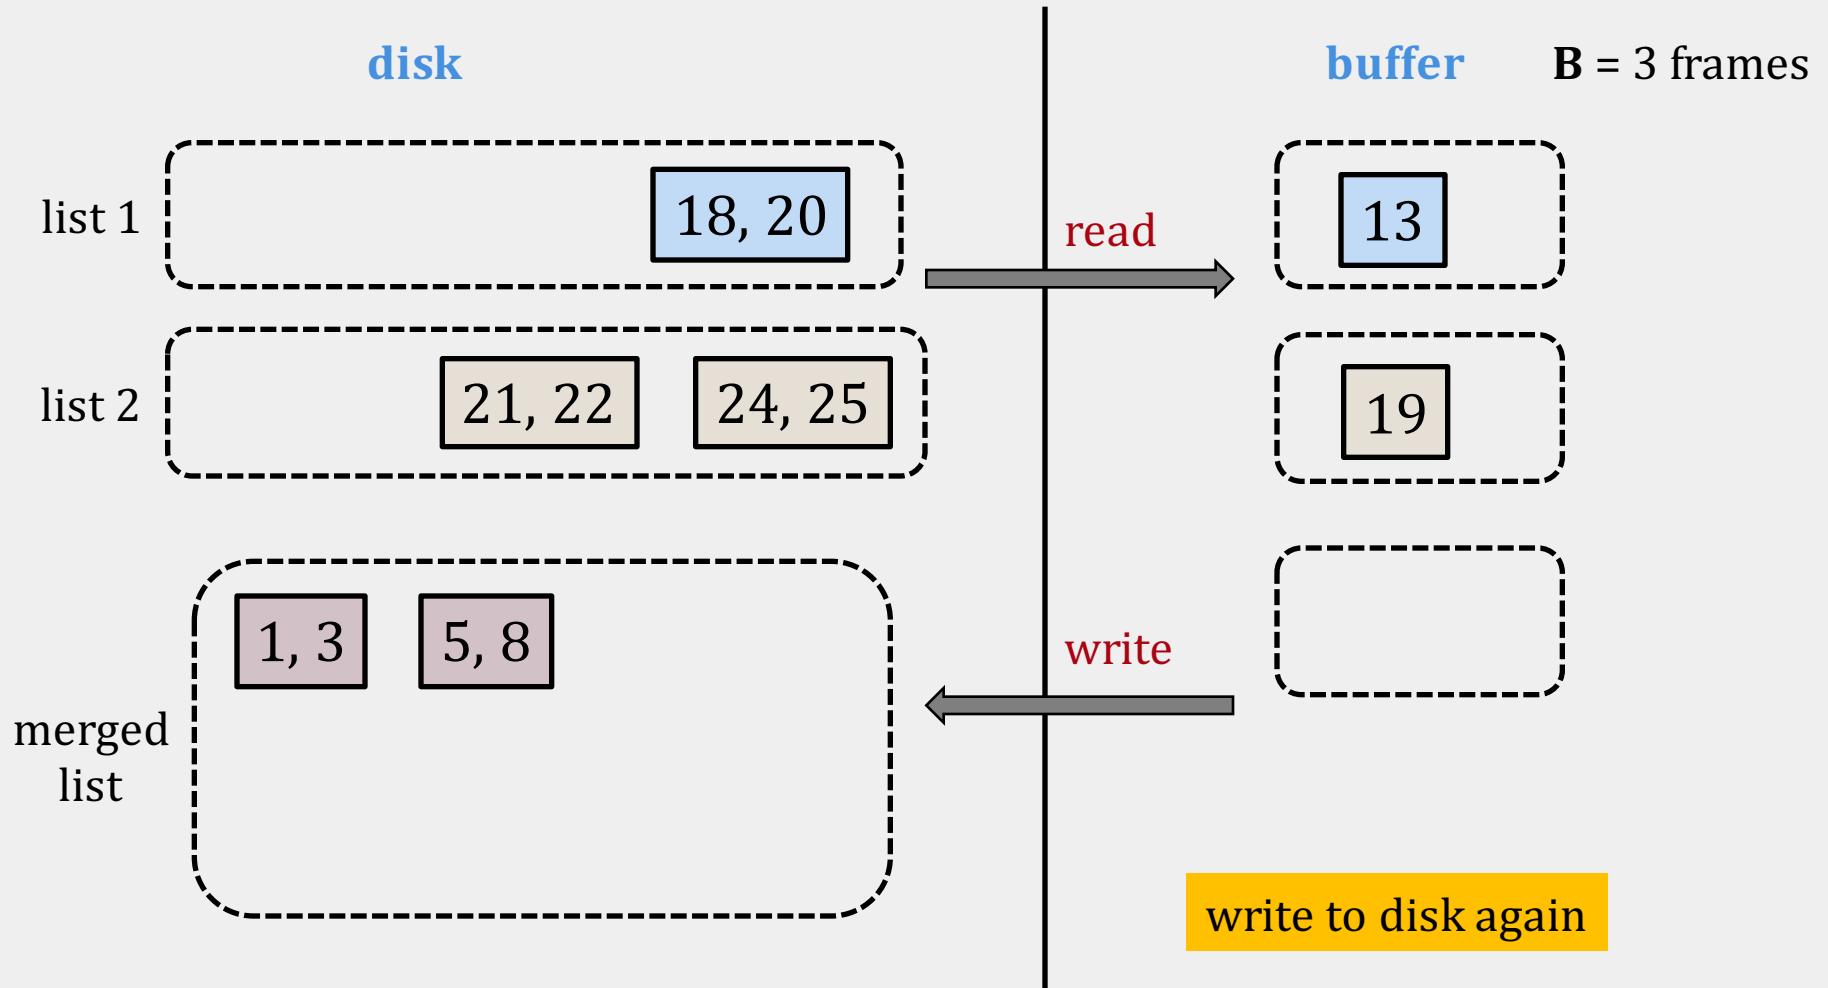
# EXTERNAL MERGE ALGORITHM



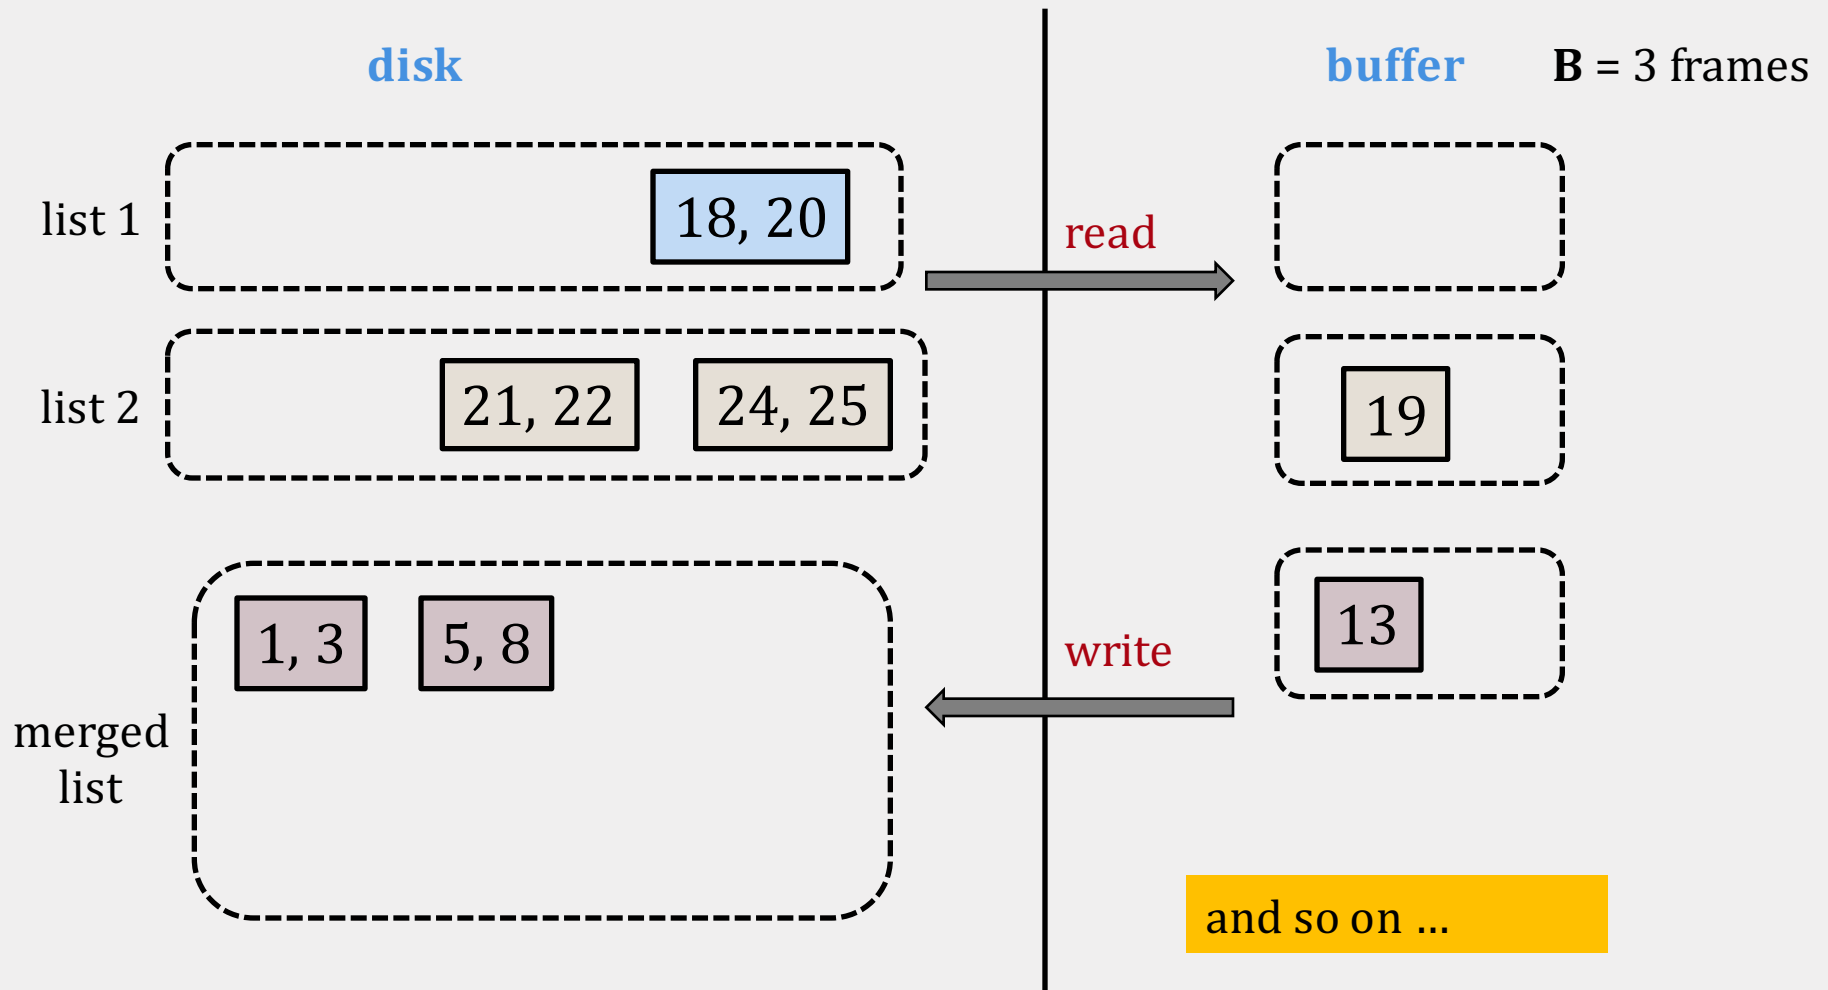
# EXTERNAL MERGE ALGORITHM



# EXTERNAL MERGE ALGORITHM



# EXTERNAL MERGE ALGORITHM



# EXTERNAL MERGE COST

We can merge 2 sorted lists of  $M$  and  $N$  pages using 3 buffer frames with

$$\text{I/O cost} = 2 (M+N)$$

When we have  $B+1$  buffer pages, we can merge  $B$  lists with the same I/O cost



---

# EXTERNAL MERGE SORT

---

---

# THE SORTING PROBLEM

---

- **B** available pages in buffer pool
- a relation **R** of size **N** pages (where **N** > **B**)

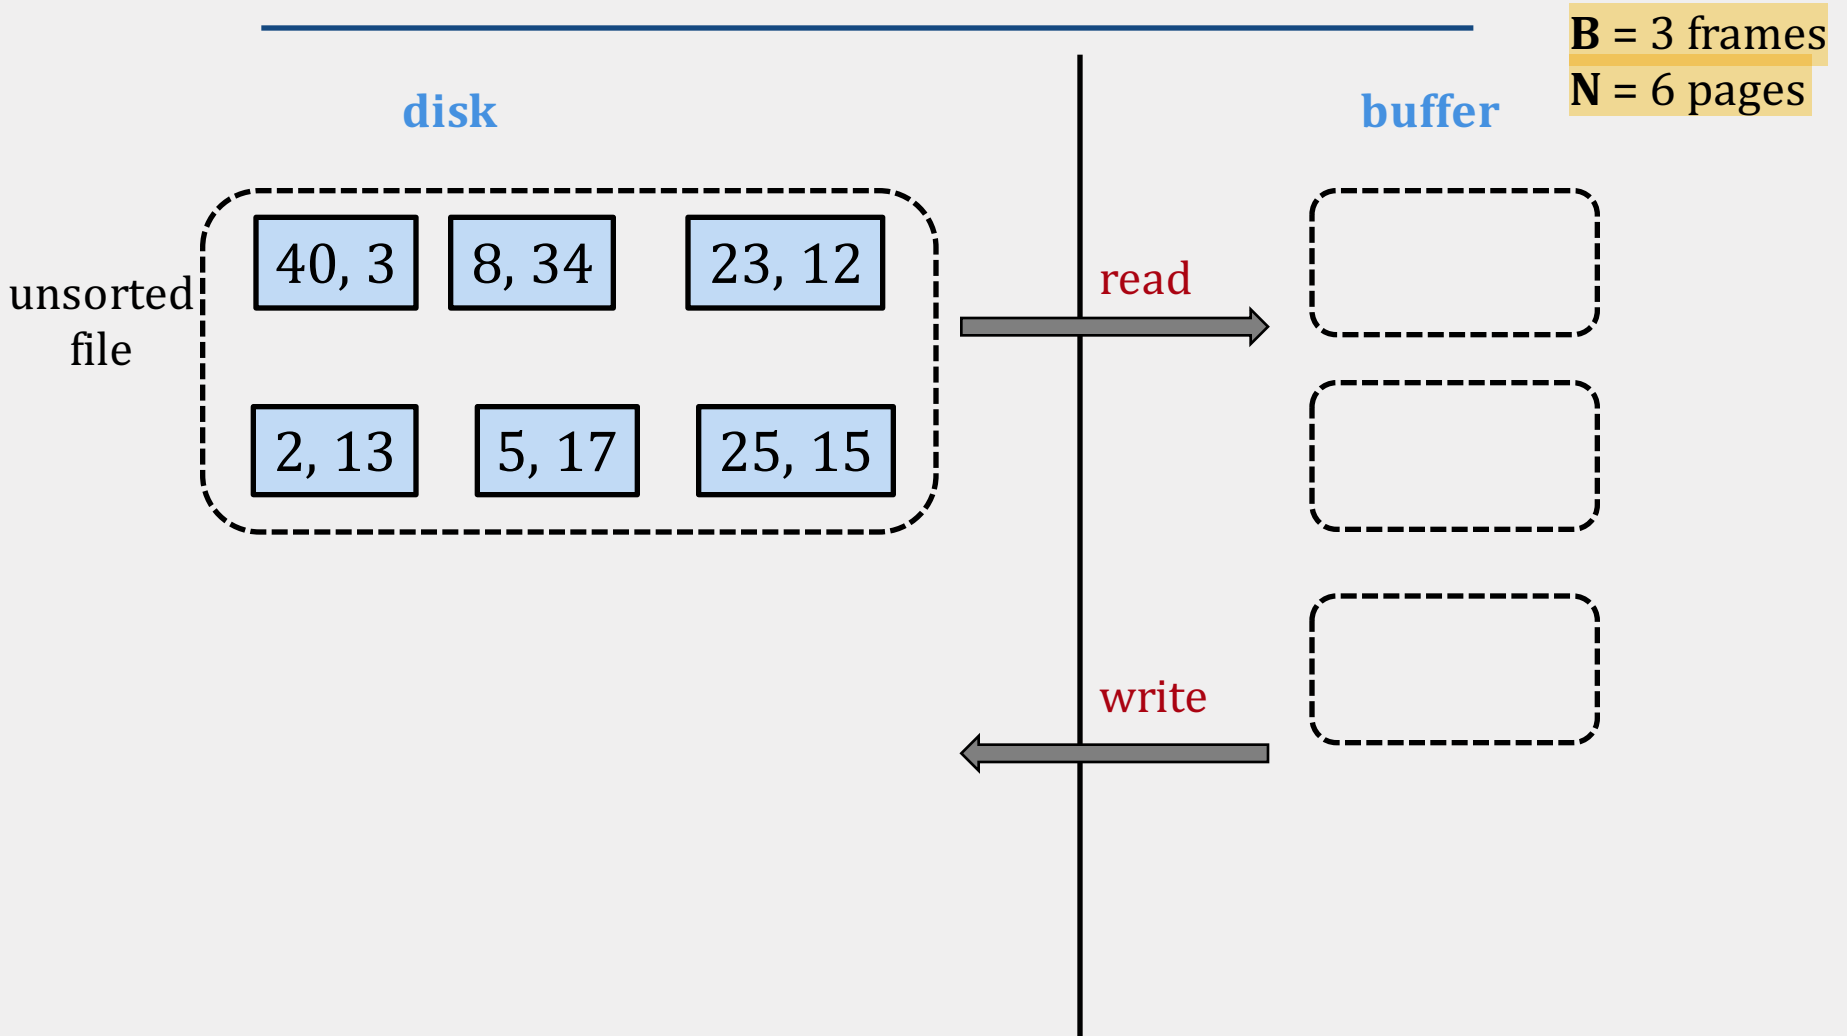
**SORTING**: output the same relation sorted on a given attribute

# KEY IDEA

---

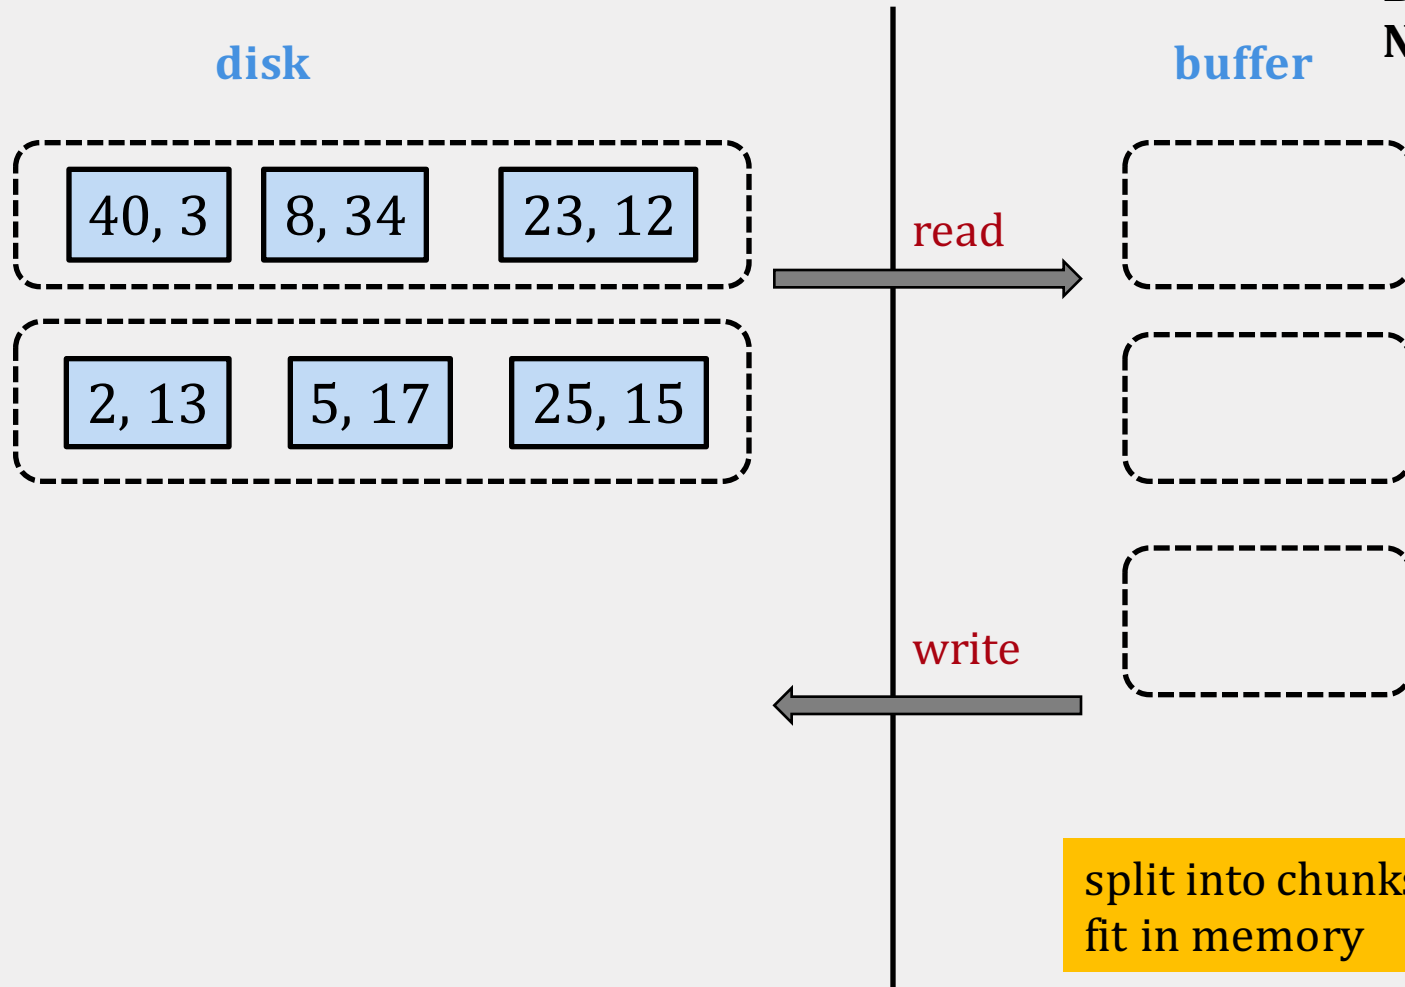
- split into chunks small enough to sort in memory (called runs)
- merge groups of runs using the external merge algorithm
- keep merging the resulting runs (each time is called a pass) until left with a single sorted file

# WARM UP: 2-WAY SORT




# WARM UP: 2-WAY SORT

B = 3 frames  
N = 6 pages



# WARM UP: 2-WAY SORT

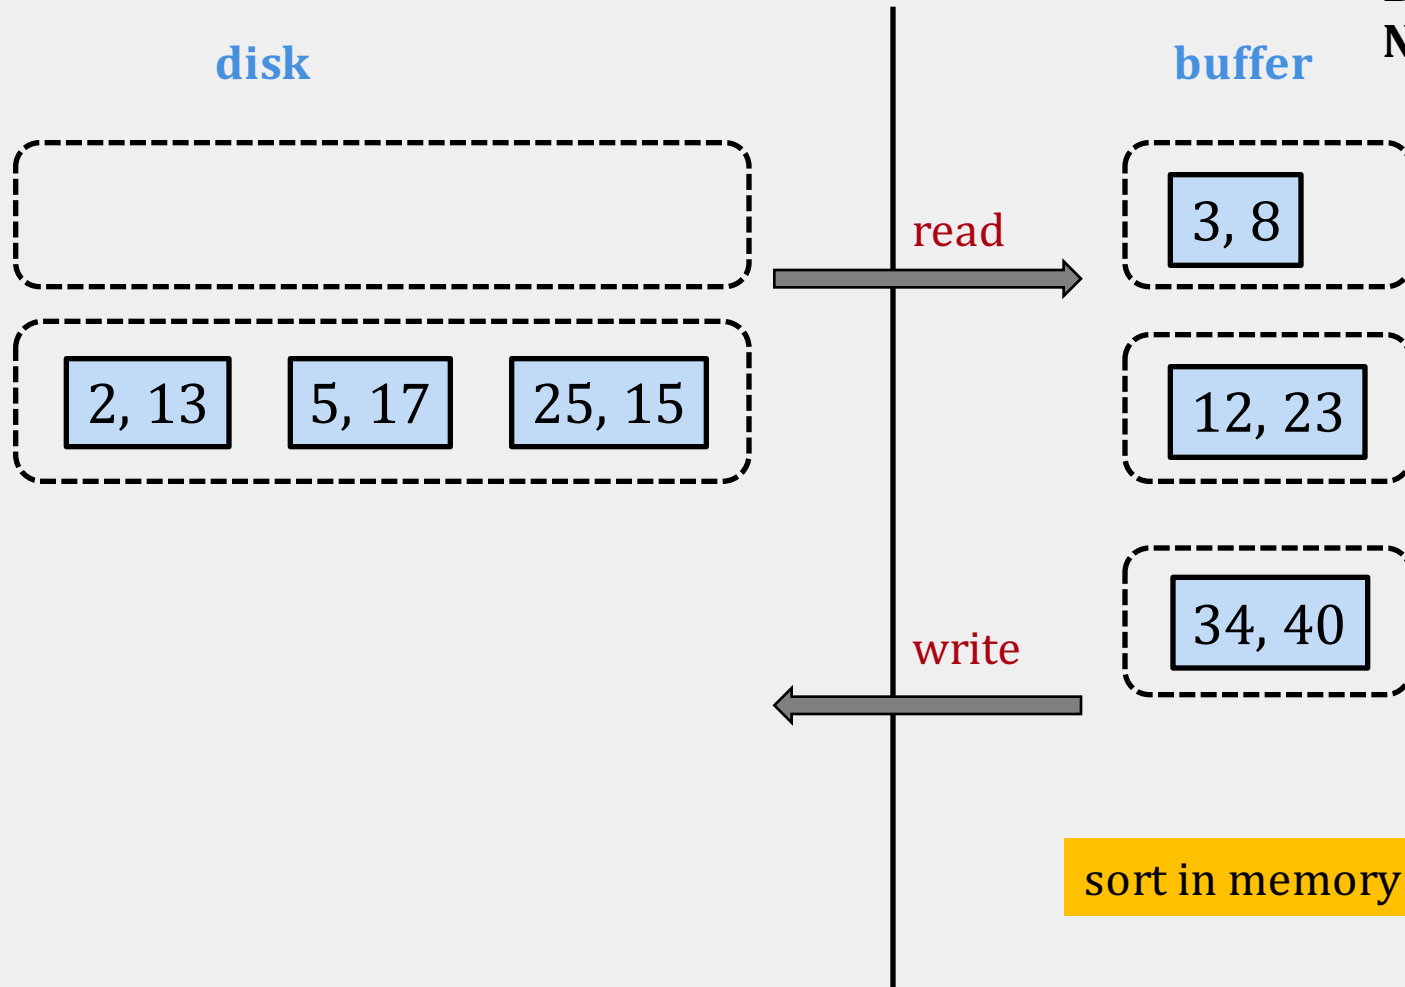
B = 3 frames  
N = 6 pages



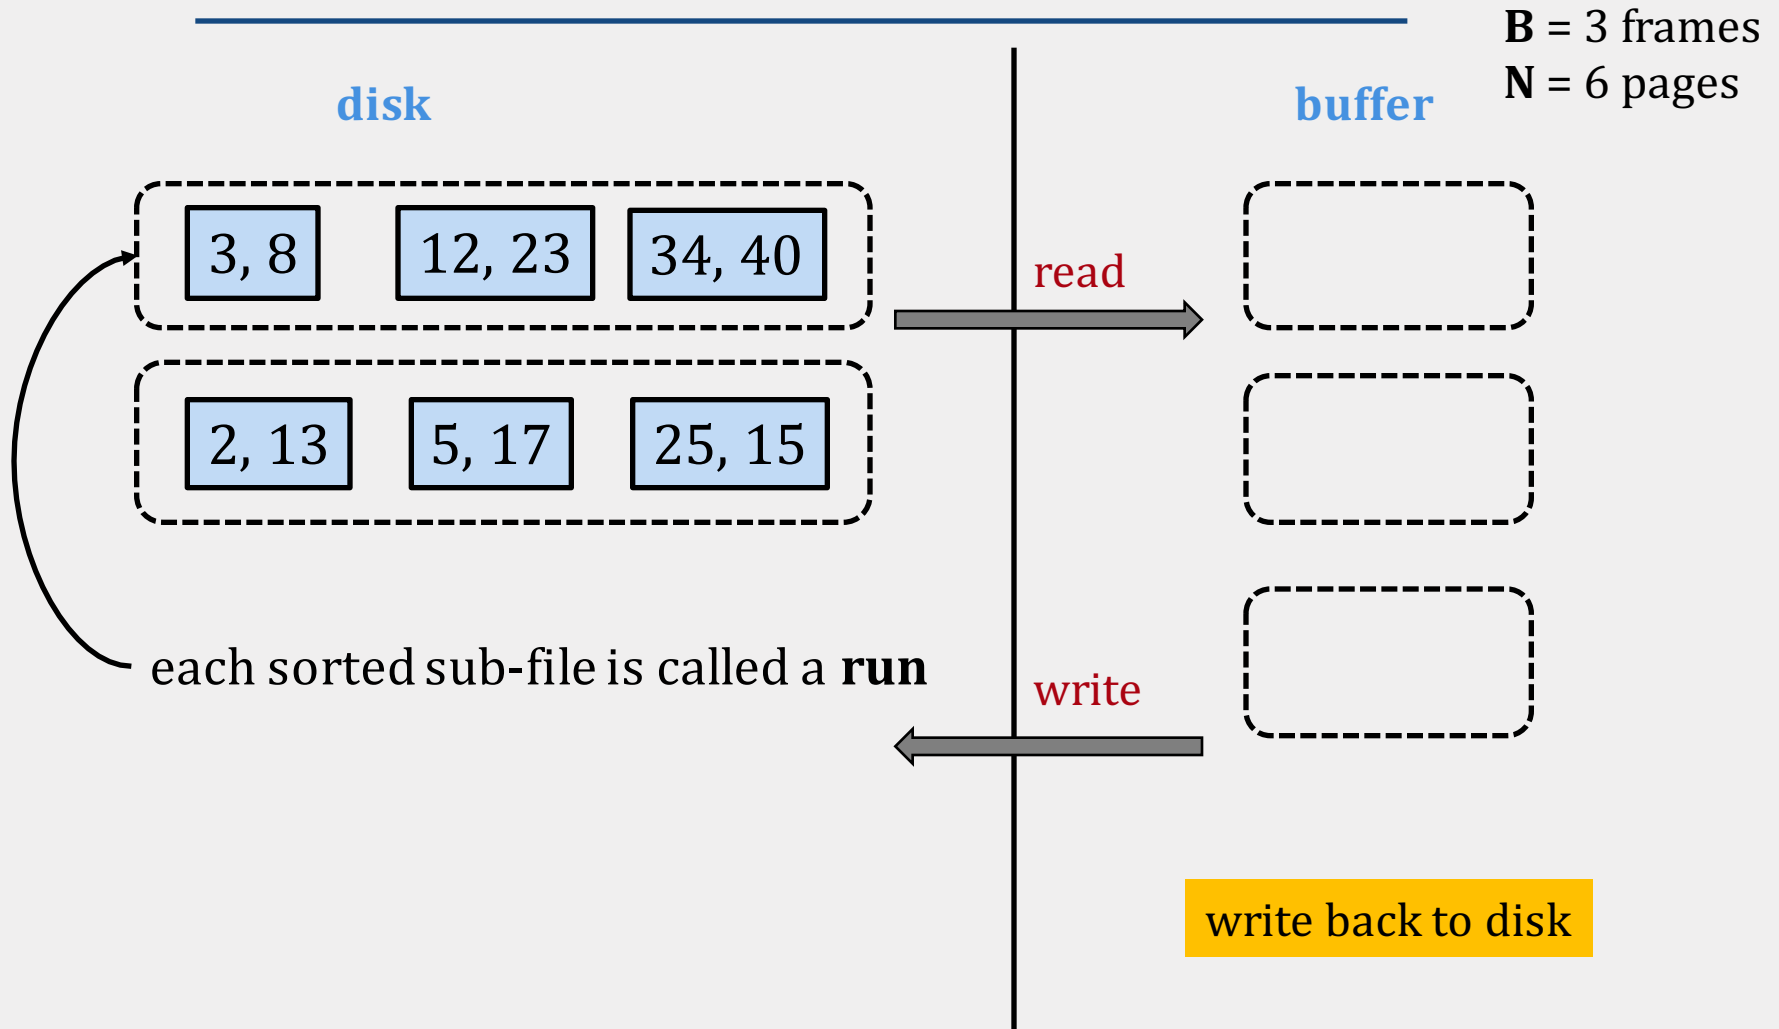
read each chunk in memory

# WARM UP: 2-WAY SORT

B = 3 frames  
N = 6 pages



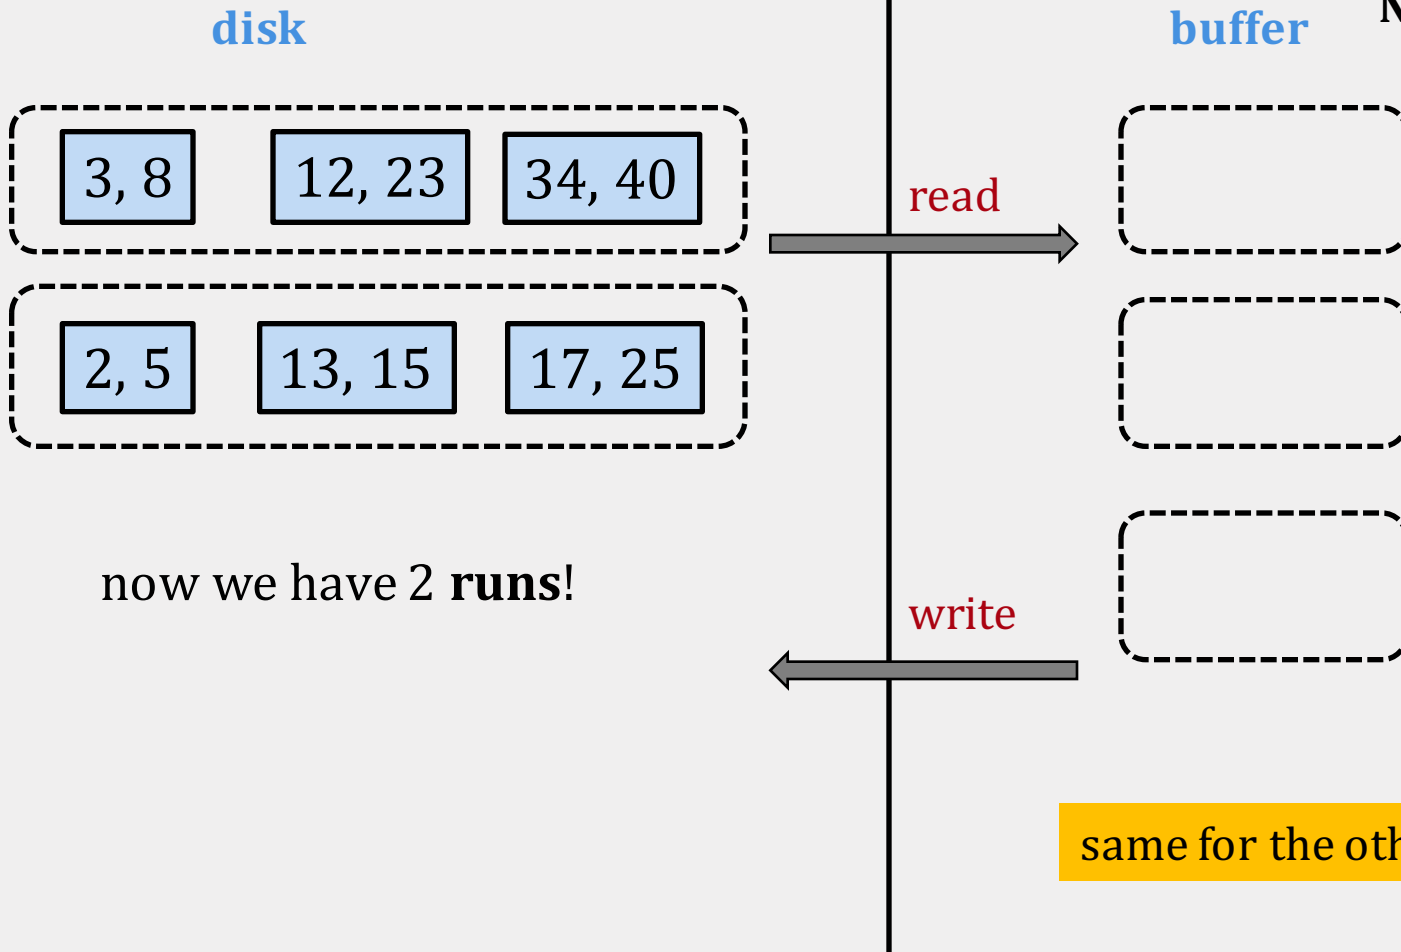
# WARM UP: 2-WAY SORT





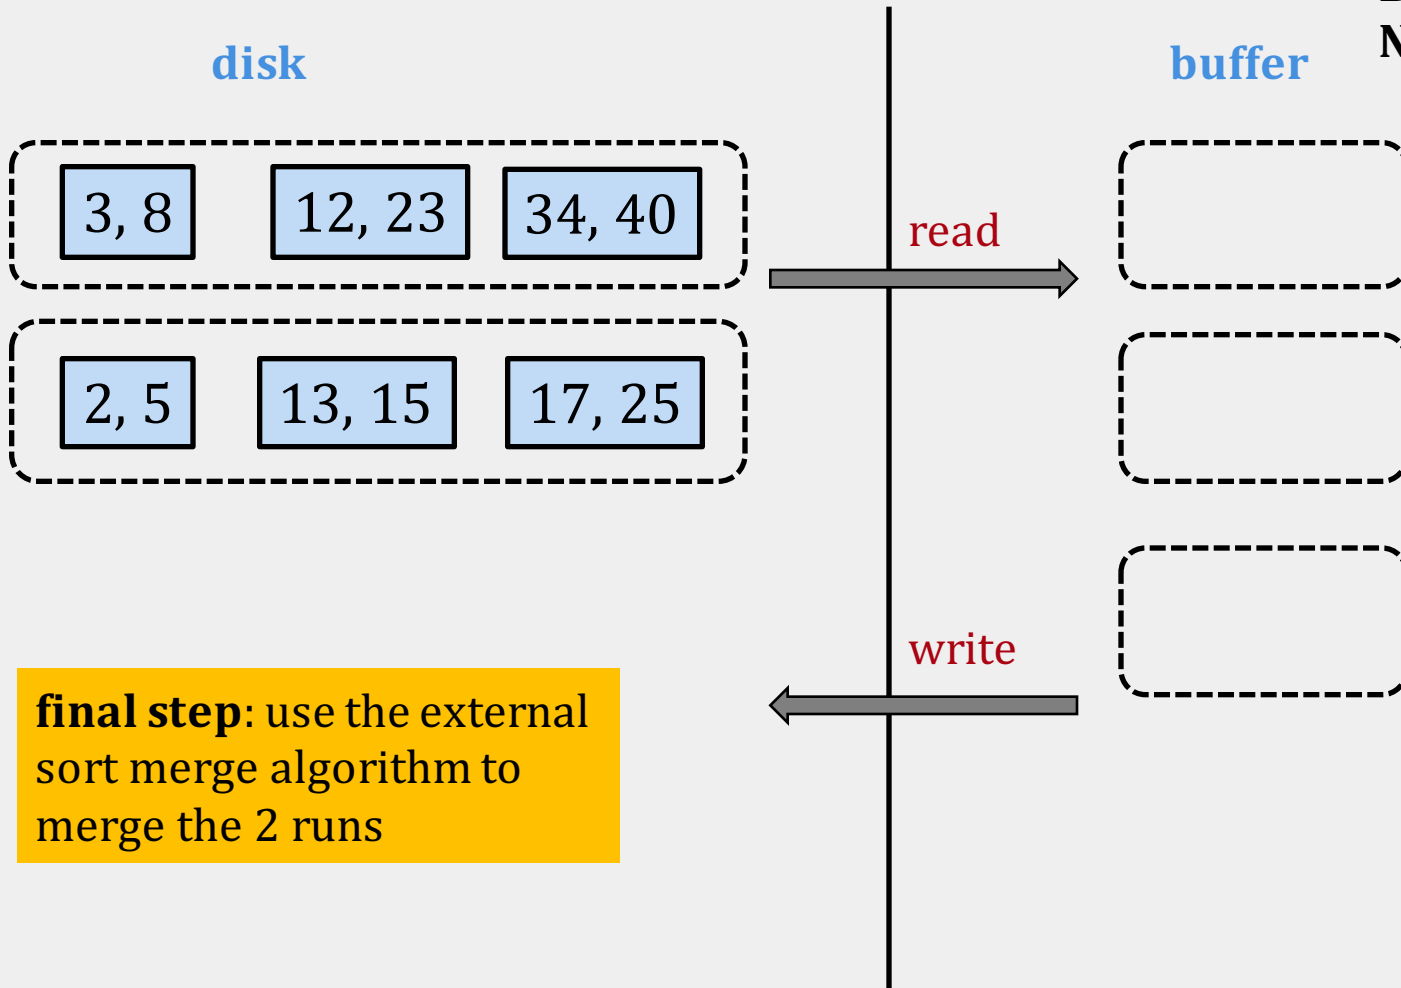
# WARM UP: 2-WAY SORT

B = 3 frames  
N = 6 pages



# WARM UP: 2-WAY SORT



$B = 3$  frames  
 $N = 6$  pages



**final step:** use the external sort merge algorithm to merge the 2 runs

# CALCULATING THE I/O COST


In our example, **B**= 3 buffer pages, **N** = 6 pages

- Pass **0**: creating the first runs
  - 1 read + 1 write for every page 
  - total cost =  $6 * (1 + 1) = 12$  I/Os
- Pass **1**: external merge sort 
  - total cost =  $2 * (3 + 3) = 12$  I/Os

So 24 I/Os in total

# I/O COST: SIMPLIFIED VERSION

Assume for now that we initially create **N runs**, each run consisting of **a single page**



**pass 0:** N runs, each 1 page



**pass 1:** merge into N/2 runs

**pass 2:** merge into N/4 runs

- We need  $\lceil \log_2 N \rceil + 1$  passes to sort the whole file
- Each pass needs  $2N$  I/Os

$$\text{total I/O cost} = 2N(\lceil \log_2 N \rceil + 1)$$



# CAN WE DO BETTER?

- The 2-way merge algorithm only uses 3 buffer pages
- But we have more available memory!

**Key idea:** use as much of the available memory as possible in every pass

- reducing the number of passes reduces I/O

# EXTERNAL SORT: I/O COST

Suppose we have  $B \geq 3$  buffer pages available

$$2N(\lceil \log_2 N \rceil + 1) \longrightarrow 2N(\lceil \log_2 \frac{N}{B} \rceil + 1) \longrightarrow 2N(\lceil \log_{B-1} \frac{N}{B} \rceil + 1)$$

- initial runs of length 1
- 3-way merge

increase the length of the initial runs to B

merge B-1 runs at a time

---

# NUMBER OF PASSES

---

<b>N</b>	<b>B=3</b>	<b>B=17</b>	<b>B=257</b>
100	7	2	1
10,000	13	4	2
1,000,000	20	5	3
10,000,000	23	6	3
100,000,000	26	7	4
1,000,000,000	30	8	4

---

# OPTIMIZING MERGE SORT

---



---

# REPLACEMENT SORT

---


- used as an alternative for the sorting in pass 0
- creates runs of *average* size  $2B$  (instead of  $B$ )

## Algorithm

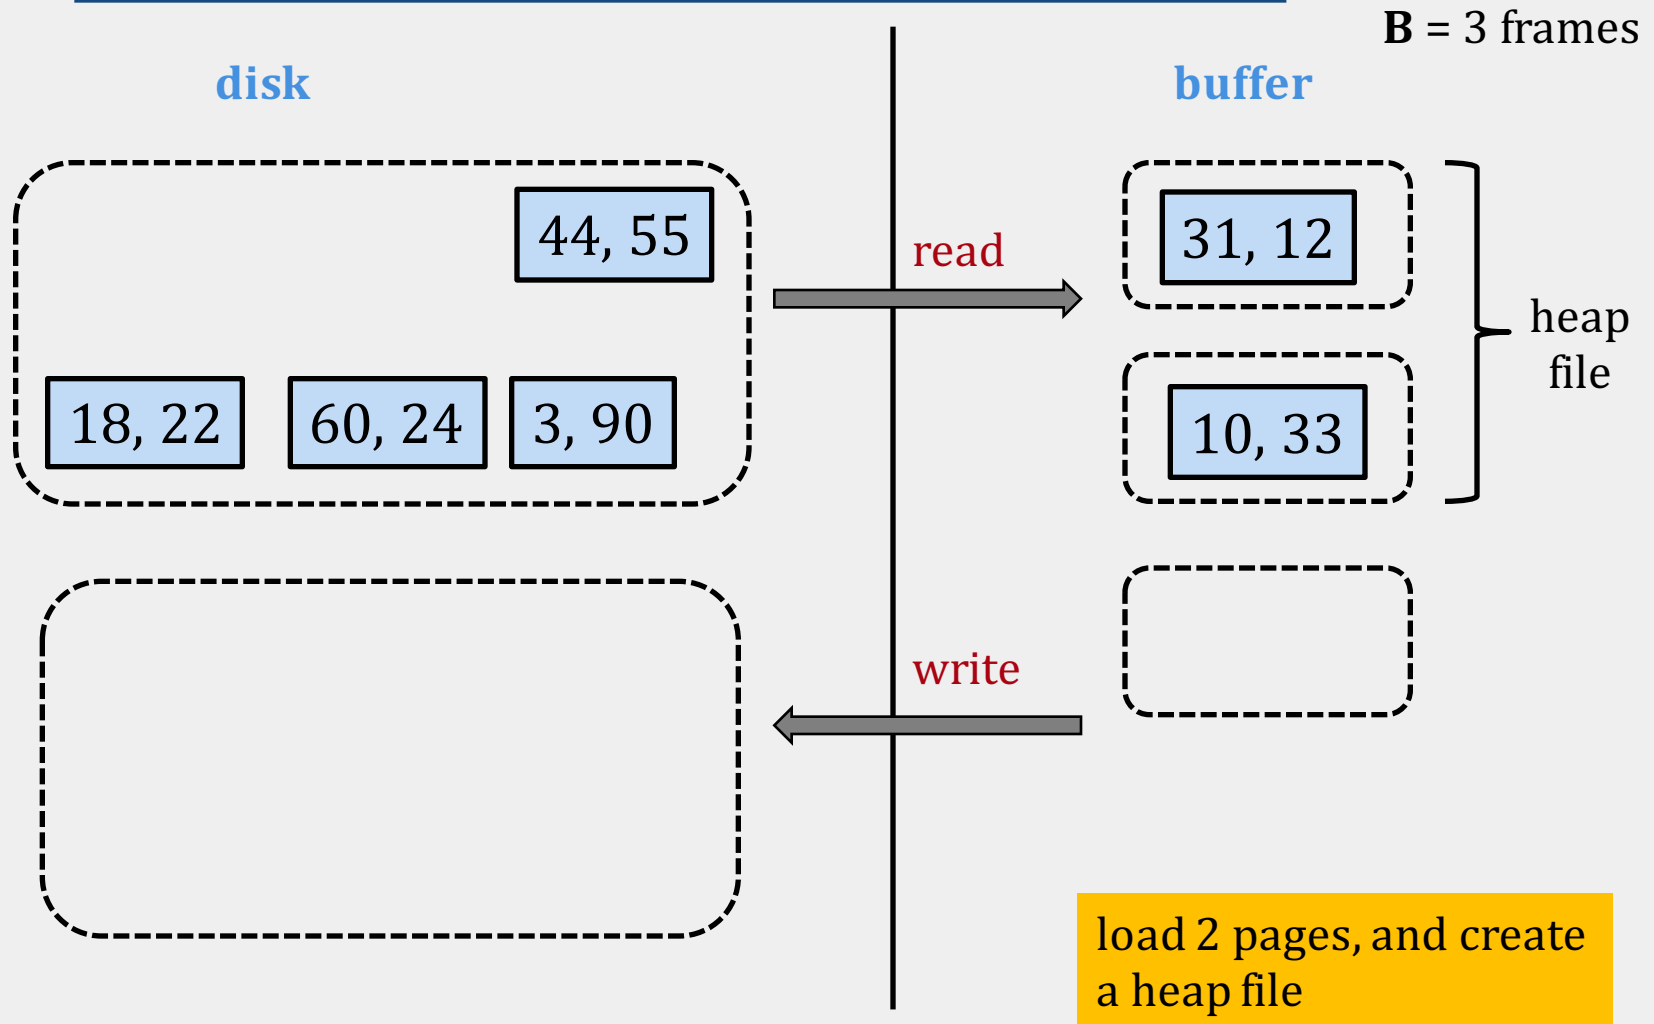
- read  $B-1$  pages in memory (keep as sorted heap)
- move smallest record (that is greater than the largest element in buffer) to output buffer
- read a new record  $r$  and insert into the sorted heap

# REPLACEMENT SORT: EXAMPLE

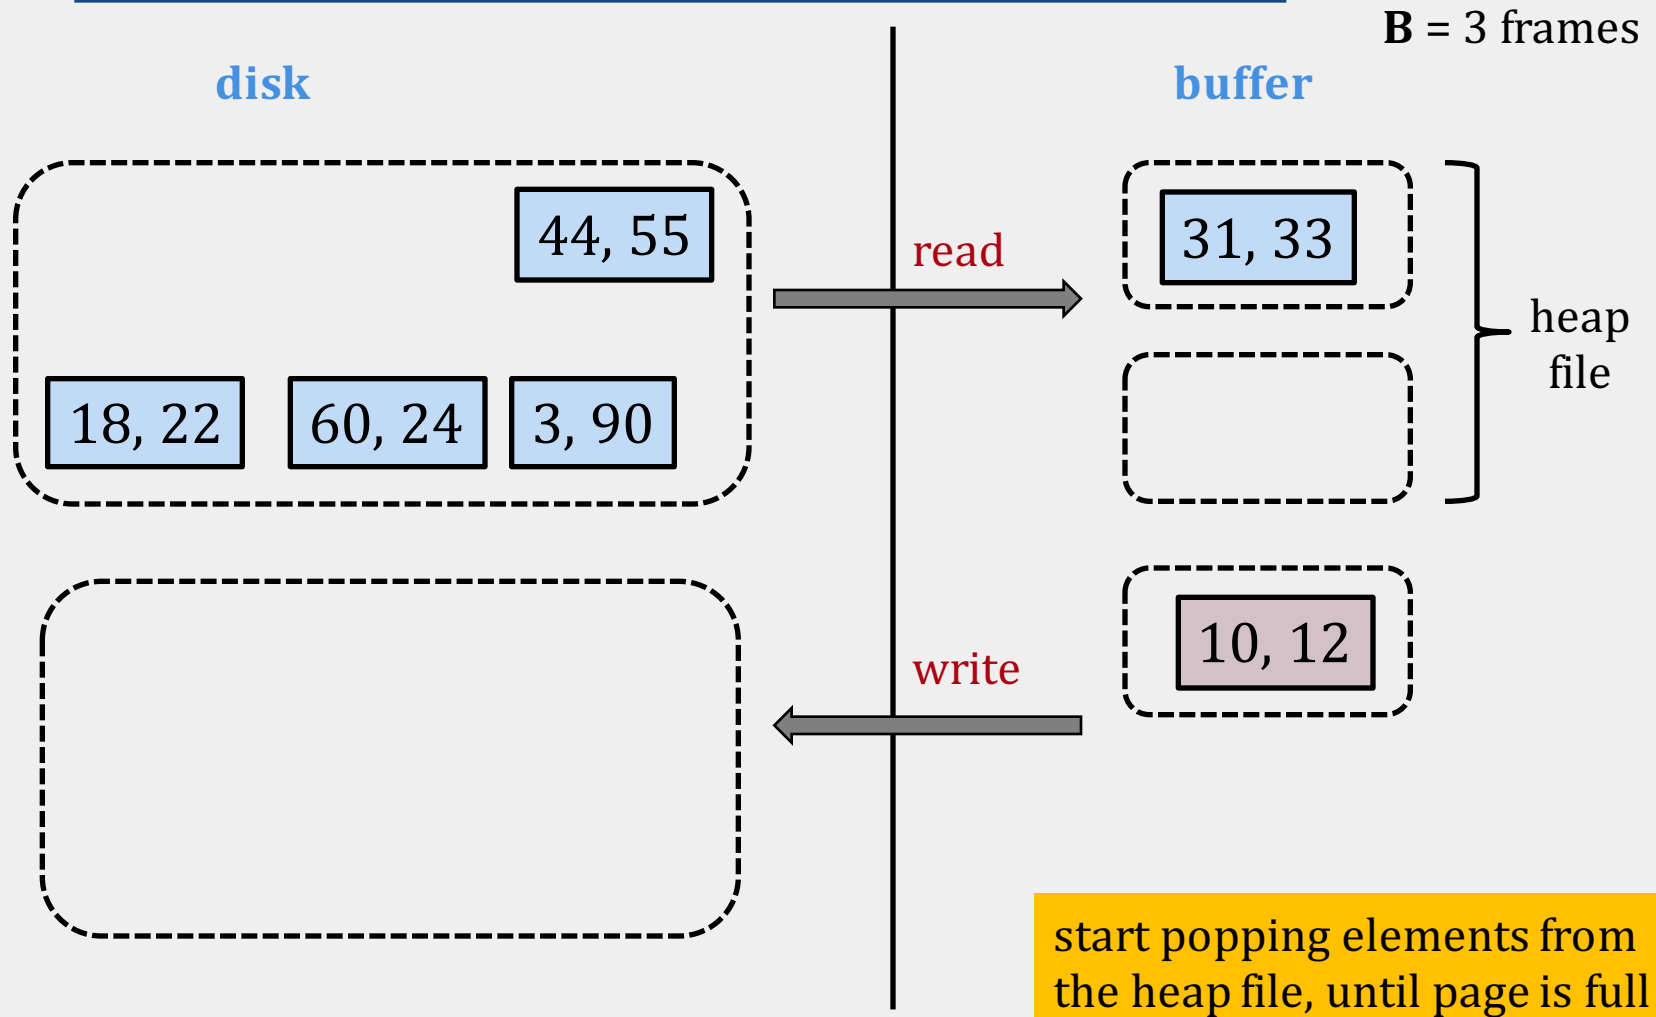
$B = 3$  frames



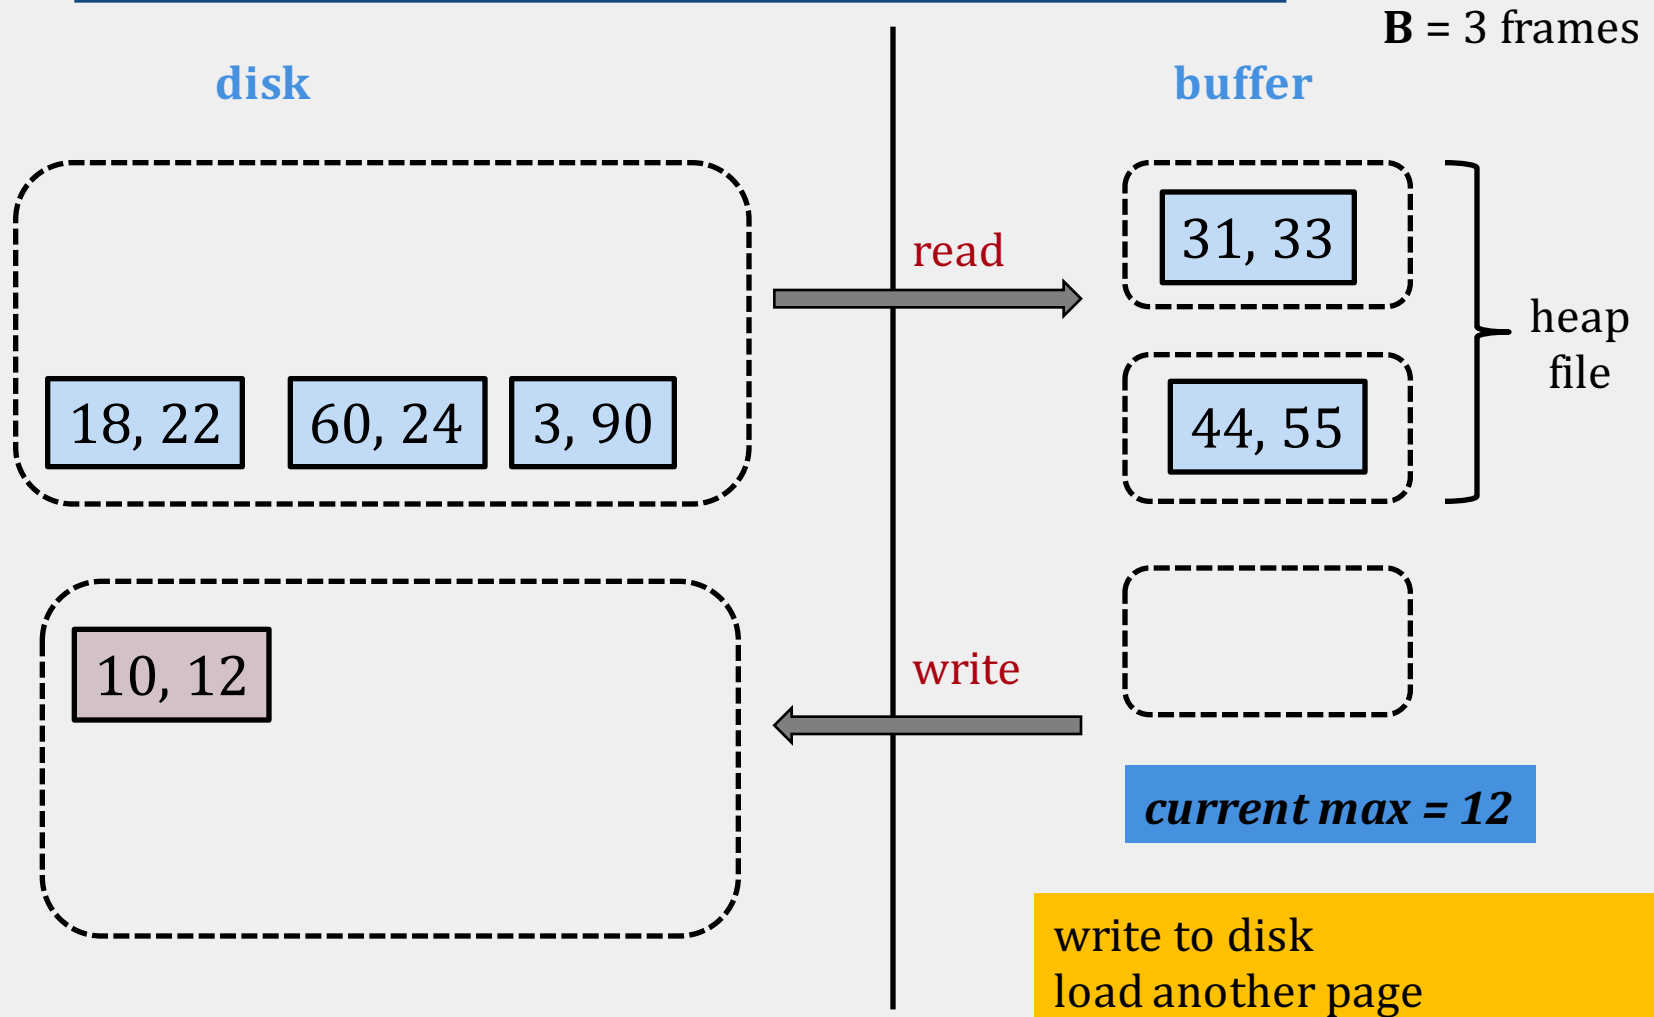
# REPLACEMENT SORT: EXAMPLE



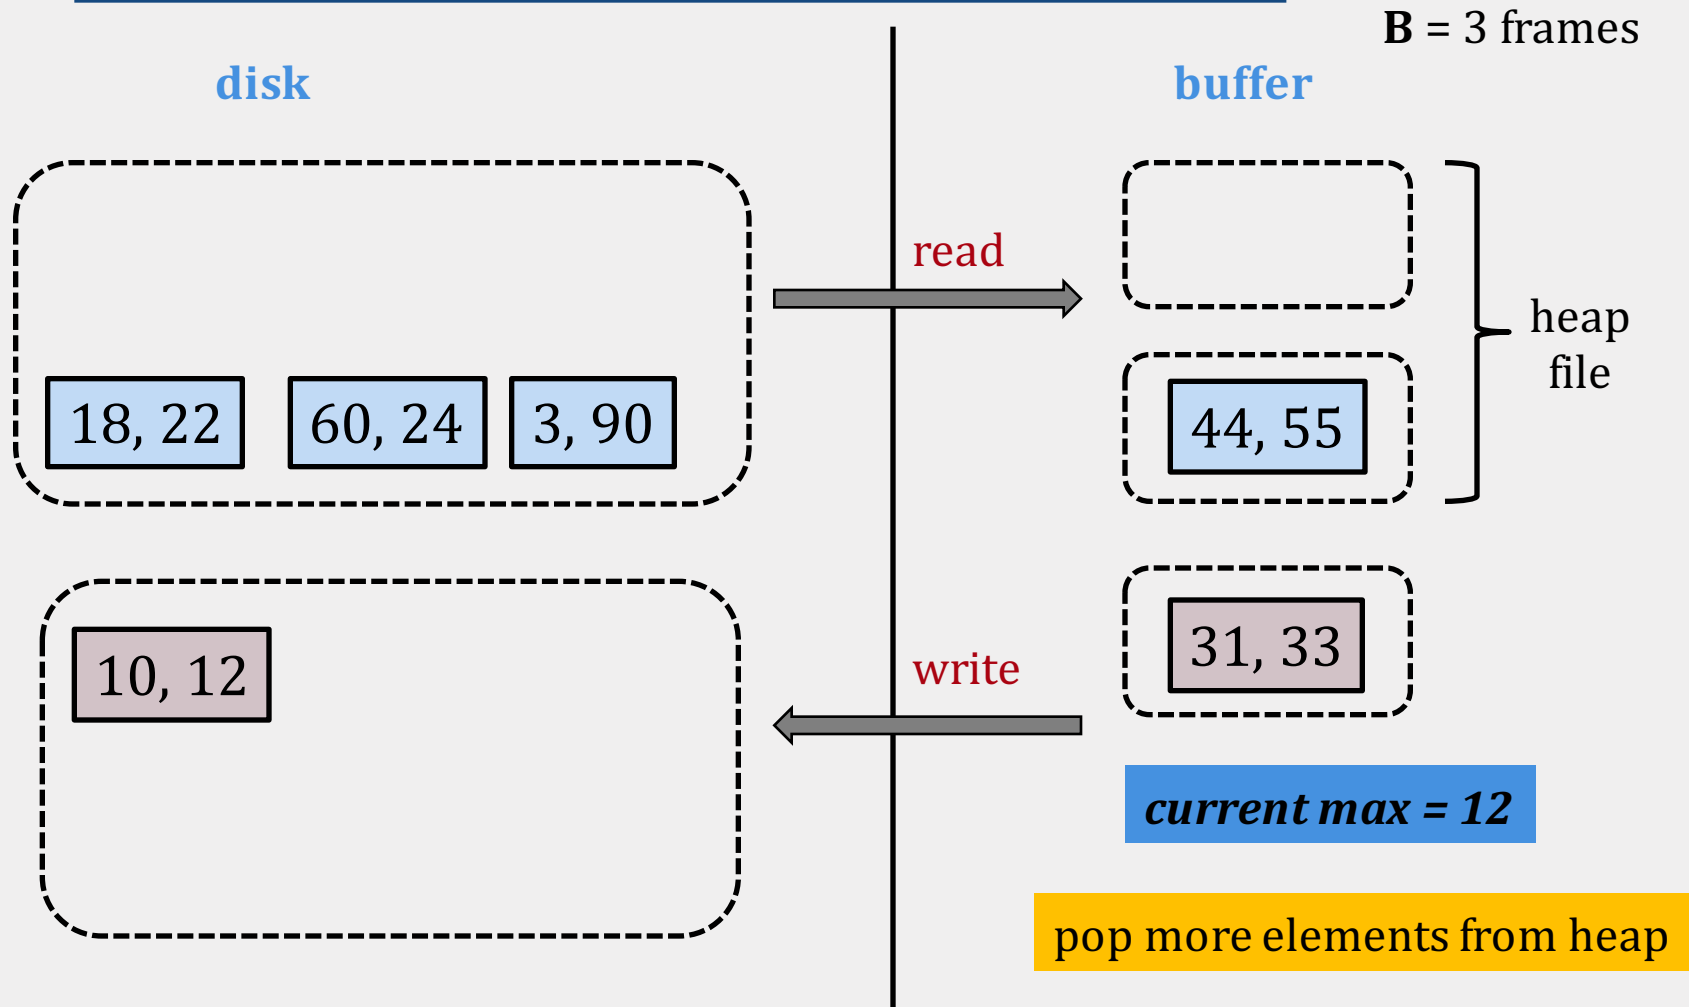
# REPLACEMENT SORT: EXAMPLE



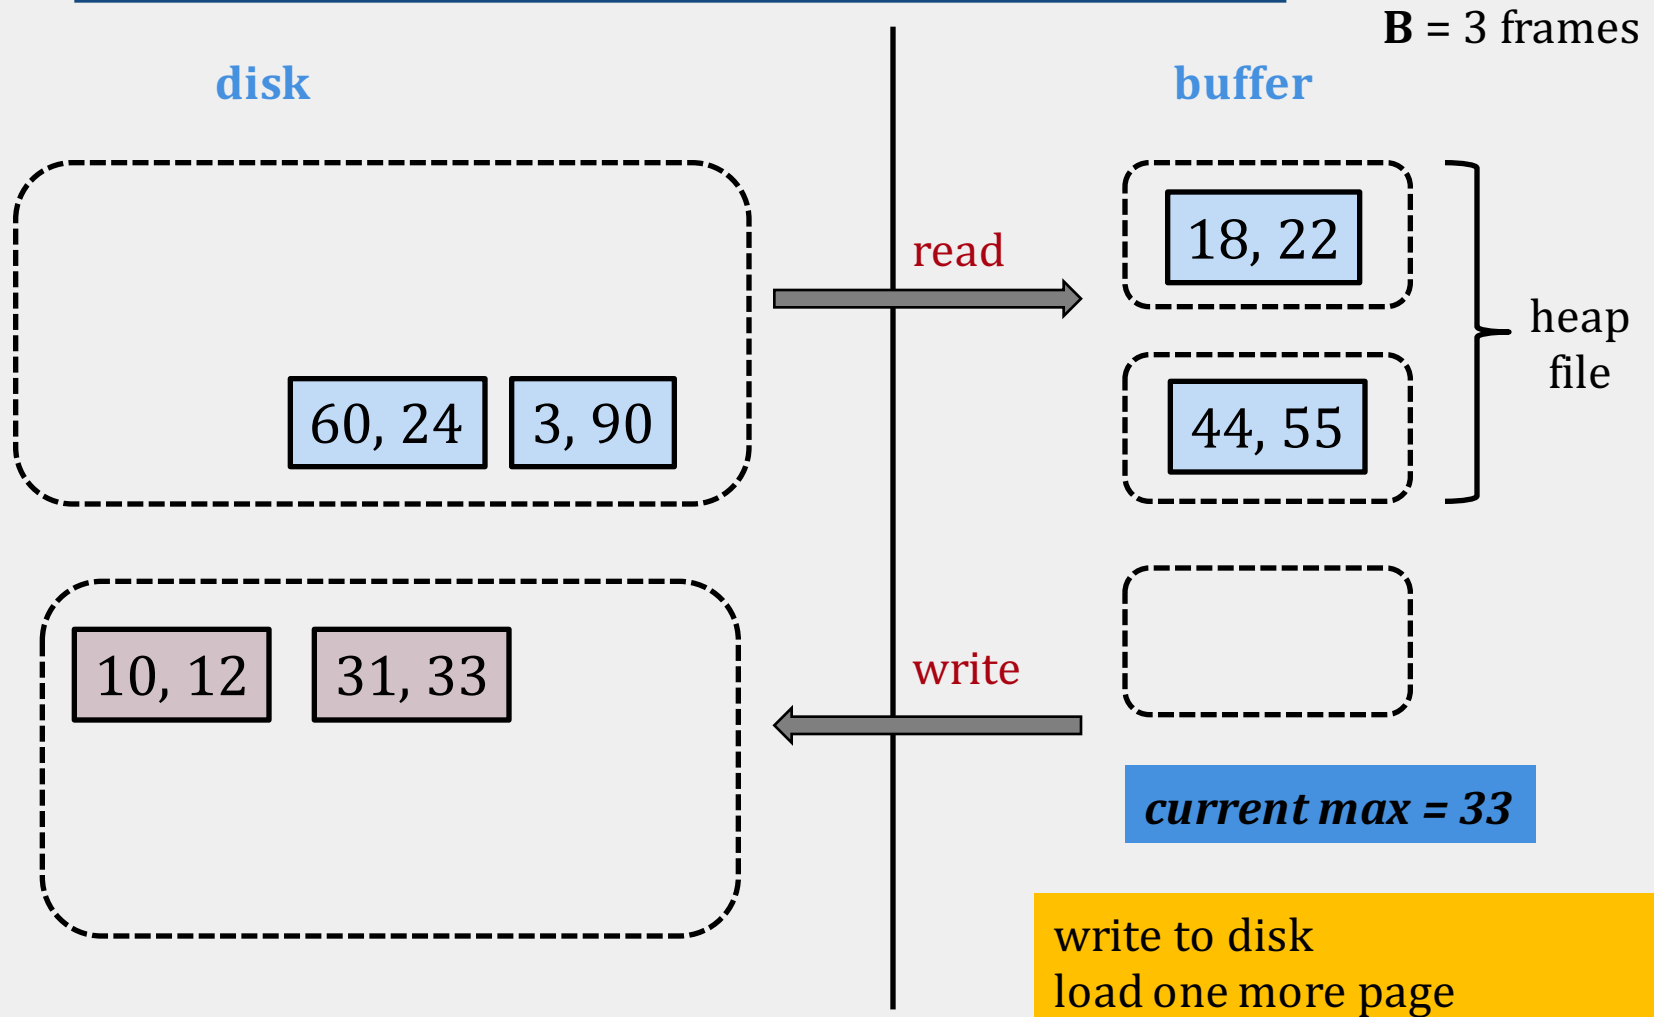
# REPLACEMENT SORT: EXAMPLE



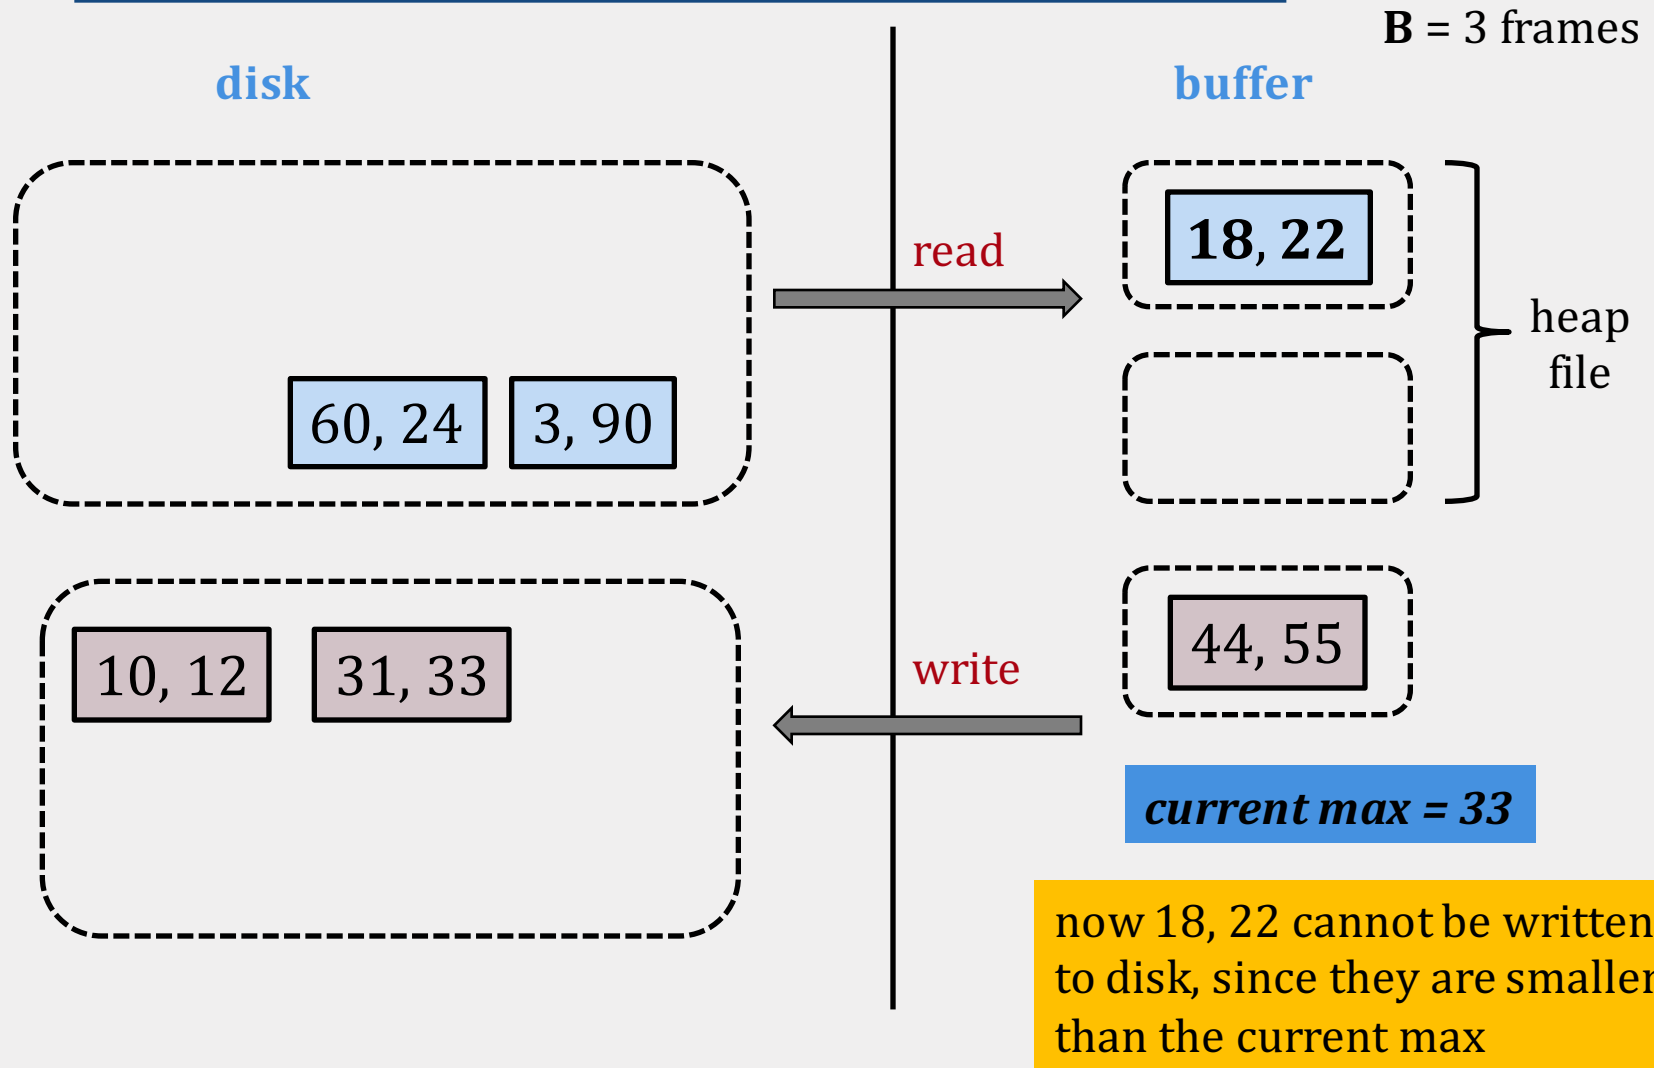
# REPLACEMENT SORT: EXAMPLE



# REPLACEMENT SORT: EXAMPLE

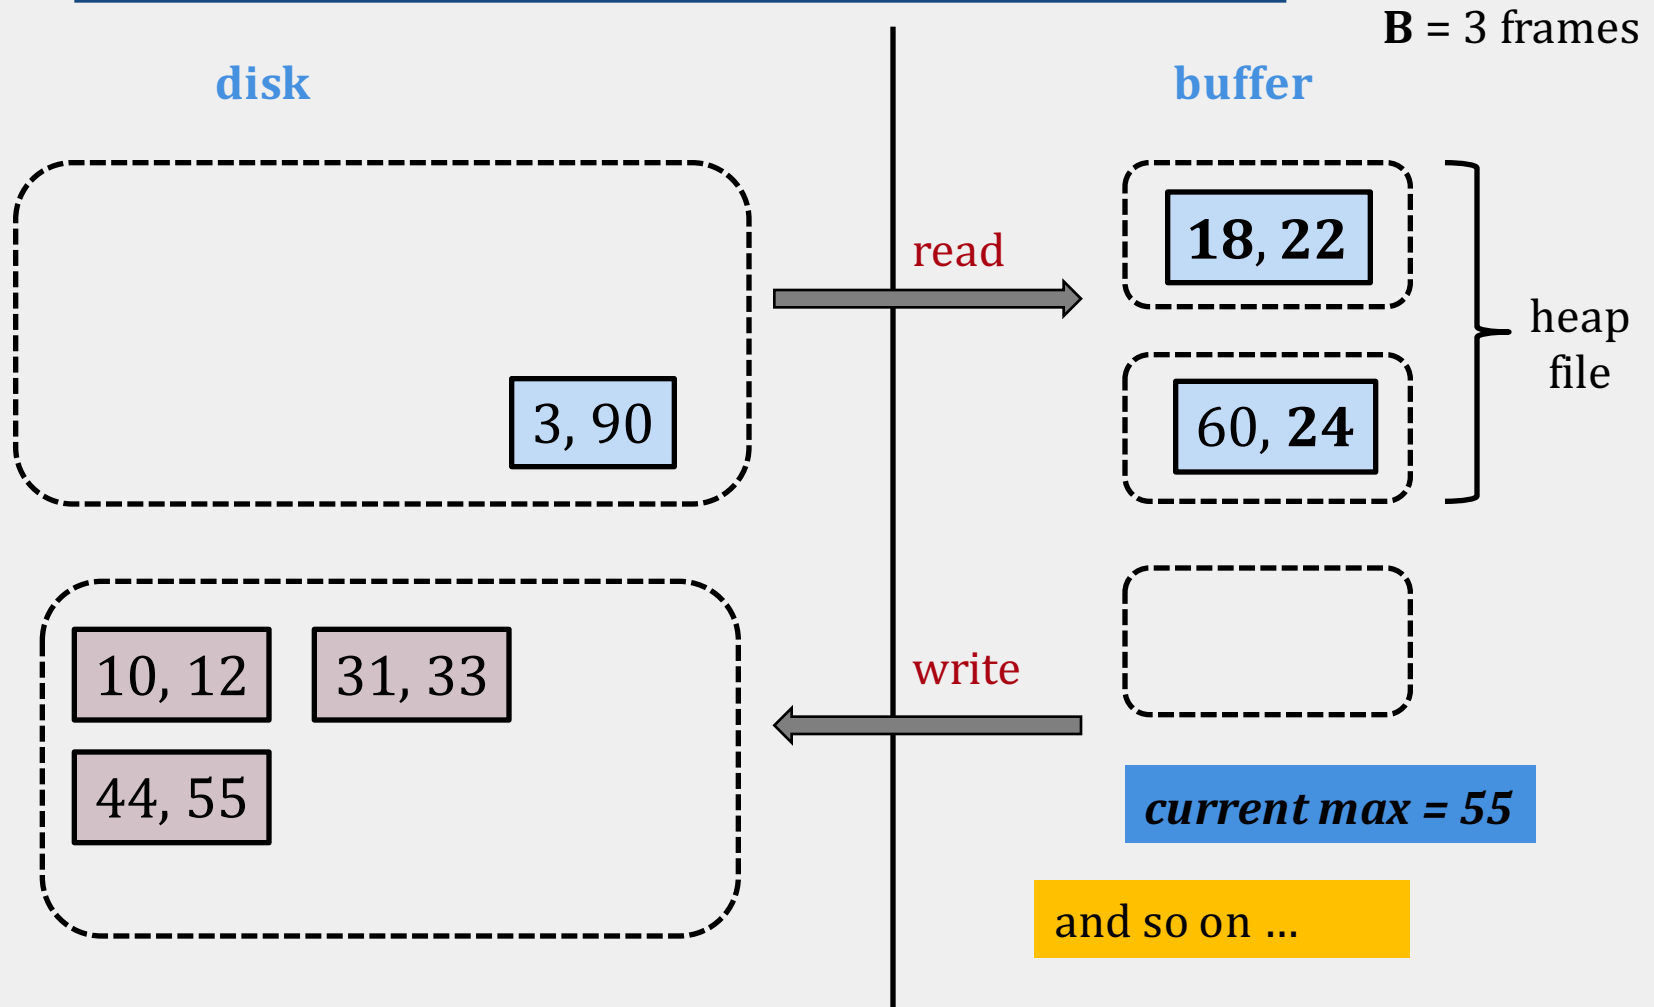


# REPLACEMENT SORT: EXAMPLE

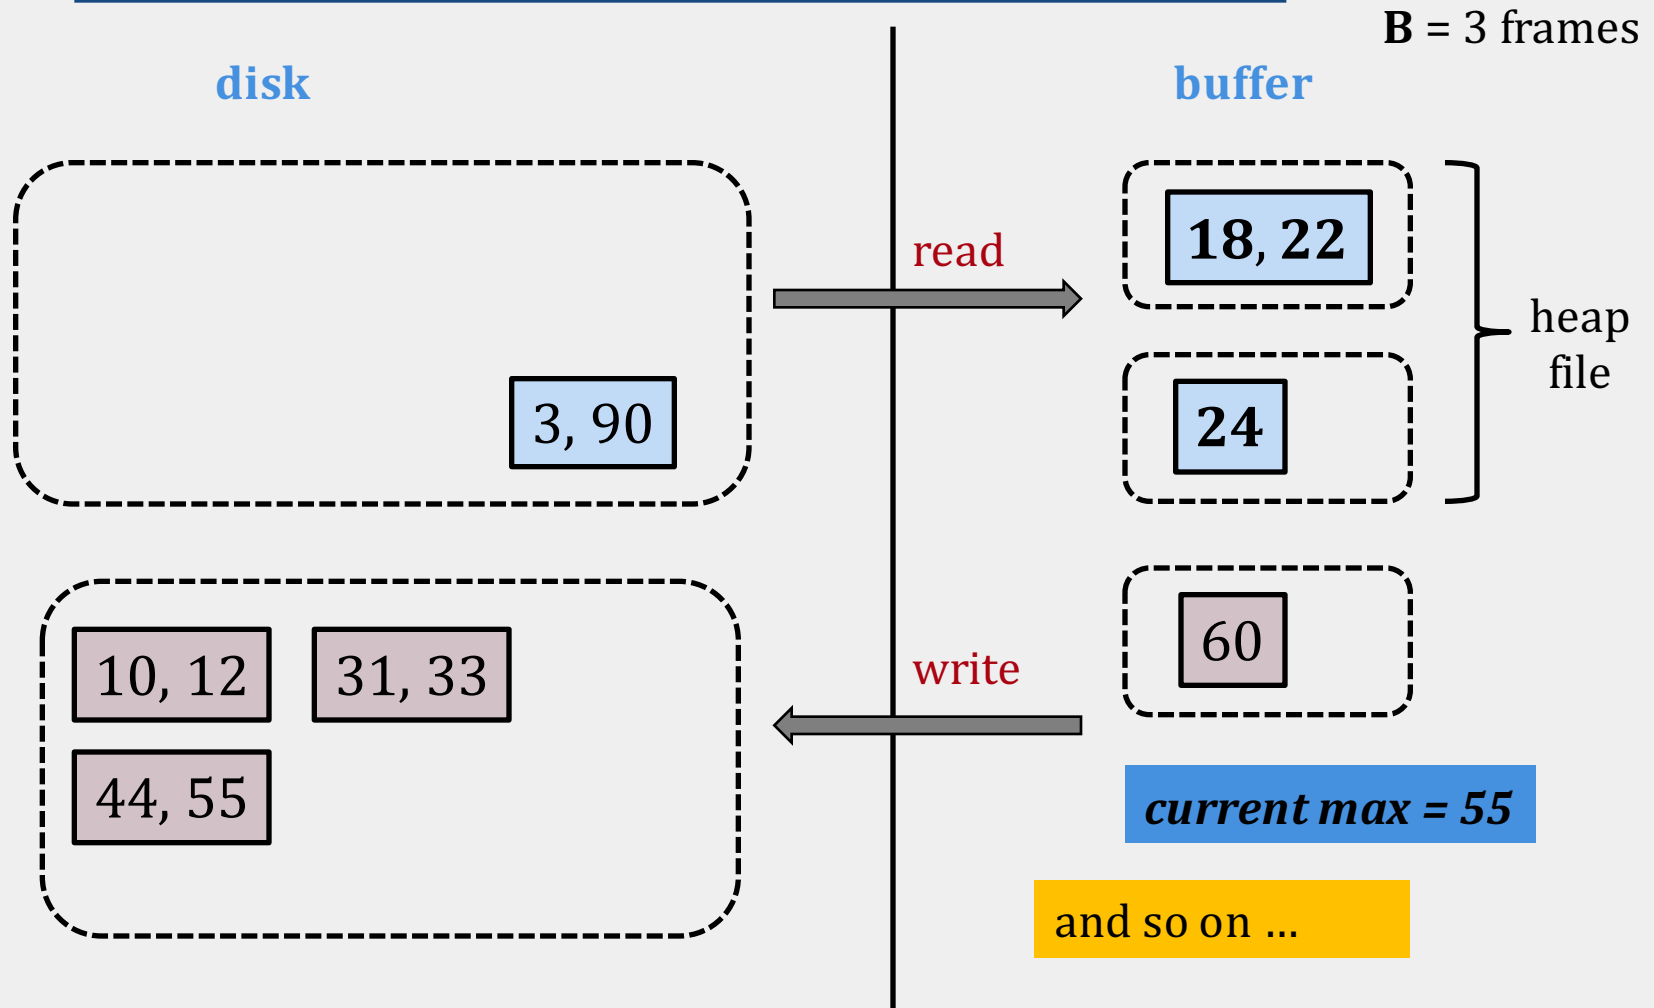




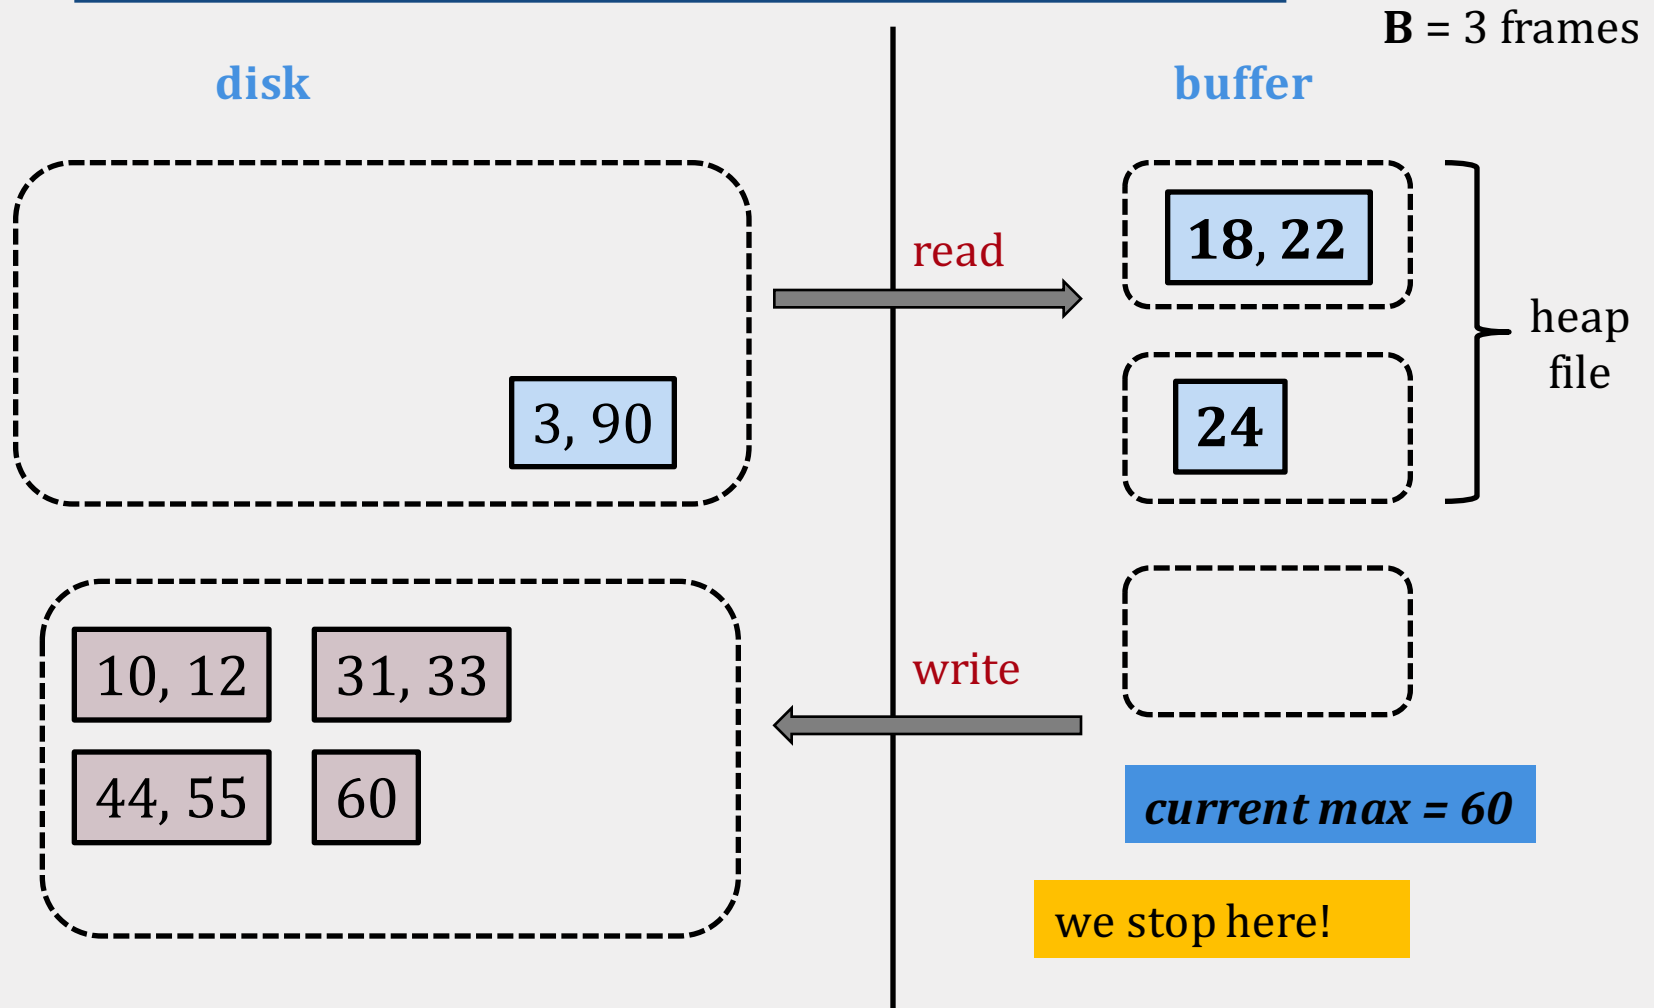
# REPLACEMENT SORT: EXAMPLE



# REPLACEMENT SORT: EXAMPLE



# REPLACEMENT SORT: EXAMPLE



# I/O COST WITH REPLACEMENT SORT

Each initial run has length  $\sim 2B$

$$\text{I/O cost} = 2N \left( \left\lceil \log_{B-1} \frac{N}{2B} \right\rceil + 1 \right)$$